

# Deep Learning with TensorFlow

Himanshu Mittal

himanshu.mittal224@gmail.com

Department of Computer Science  
Jaypee Institute of Information Technology, Noida

March 24, 2018

# Outline

## Introduction to Deep Learning

Deep Learning Success

Why Deep Learning and why now?

## TensorFlow

Introduction to TensorFlow

TensorFlow Example

Linear Regression & Gradient Descent

## Fundamentals of Deep Learning

Deep Neural Network

Artificial Neural Network

## Computer Vision & Machine Learning

Computer Vision in Machine Learning

## Artificial Neural Networks (ANN)

ANN for Computer Vision

Neural Network for Image Classification

## Image Classification using CNN

Introduction to CNN

Case Studies of CNN

Applications

Open Problems

CNN using Tensorflow

# Outline

## Introduction to Deep Learning

### Deep Learning Success

Why Deep Learning and why now?

## TensorFlow

Introduction to TensorFlow

TensorFlow Example

Linear Regression & Gradient Descent

## Fundamentals of Deep Learning

Deep Neural Network

Artificial Neural Network

## Computer Vision & Machine Learning

Computer Vision in Machine Learning

## Artificial Neural Networks (ANN)

ANN for Computer Vision

Neural Network for Image Classification

## Image Classification using CNN

Introduction to CNN

Case Studies of CNN

Applications

Open Problems

CNN using Tensorflow

# 2016-17: Year of Deep Learning



## 2016: The Year That Deep Learning Took Over the Internet

WIRED - Dec 25, 2016

The project is still in the early stages, but it hints at the widespread impact of **deep learning** over past year. In 2016, this very old but newly ...

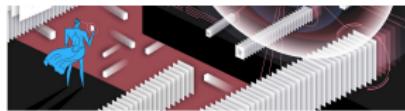


Illustration by Samathree Studio, Chicago with Ane Wang

### Deep learning takes on physics

12/08/16 | By Molly Dimond

Can the same type of technology Facebook uses to recognize faces also recognize particles?

## BuzzFeedNEWS

News Videos Quizzes Tasty DIY More ▾ Get Our News

BROKE



### This Is Why A Computer Winning At Go Is Such A Big Deal

People didn't think this would happen for at least 10 years; it's a sign of how far artificial intelligence has come.

posted on Mar 14, 2016 at 8:45 a.m.



### A Primer On Deep Learning

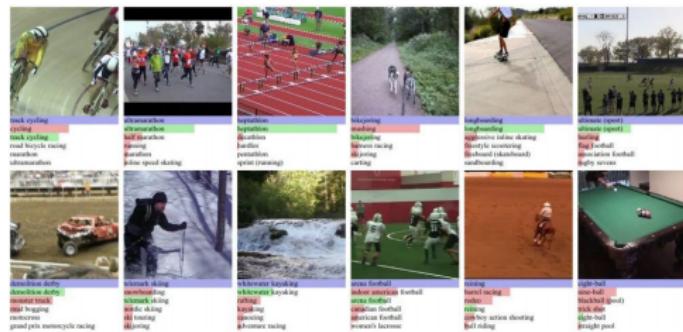
Forbes - 18-Dec-2017

Fortunately, after a half-century of research, that gap between our expectations and reality is finally closing, thanks to **deep learning**, a more advanced type of machine learning that's capable of generating human-like insight. Here, I will give an introduction to **deep learning** and explore the potential of this ...

# Deep Learning Success

- ▶ Image Classification
- ▶ Machine Translation
- ▶ Speech Recognition
- ▶ Game Playing

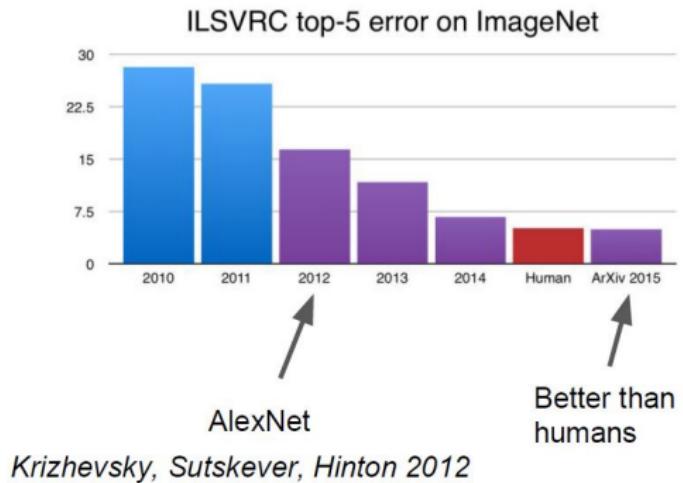
... and many, many more



# Deep Learning Success

- ▶ **Image Classification**
- ▶ Machine Translation
- ▶ Speech Recognition
- ▶ Game Playing

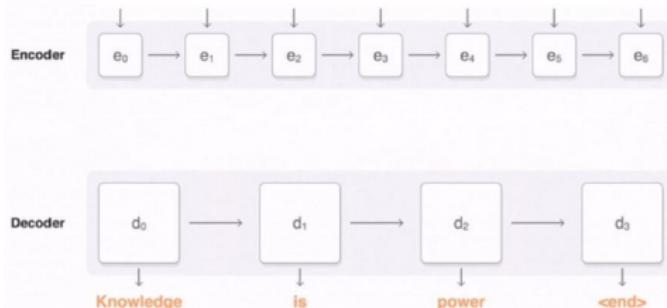
... and many, many more



# Deep Learning Success

- ▶ Image Classification
- ▶ Machine Translation
- ▶ Speech Recognition
- ▶ Game Playing

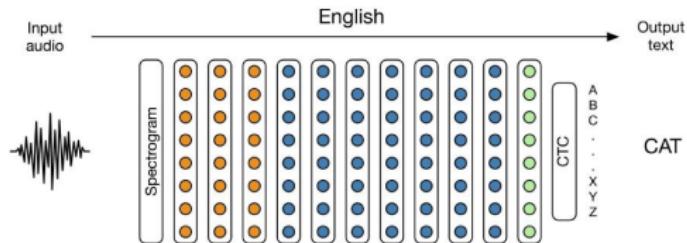
... and many, many more



# Deep Learning Success

- ▶ Image Classification
- ▶ Machine Translation
- ▶ Speech Recognition
- ▶ Game Playing

... and many, many more



## Deep Learning Success

- ▶ Image Classification
- ▶ Machine Translation
- ▶ Speech Recognition
- ▶ Game Playing

... and many, many more



# Outline

## Introduction to Deep Learning

  Deep Learning Success

  Why Deep Learning and why now?

## TensorFlow

  Introduction to TensorFlow

  TensorFlow Example

  Linear Regression & Gradient Descent

## Fundamentals of Deep Learning

  Deep Neural Network

  Artificial Neural Network

## Computer Vision & Machine Learning

  Computer Vision in Machine Learning

## Artificial Neural Networks (ANN)

  ANN for Computer Vision

  Neural Network for Image Classification

## Image Classification using CNN

  Introduction to CNN

  Case Studies of CNN

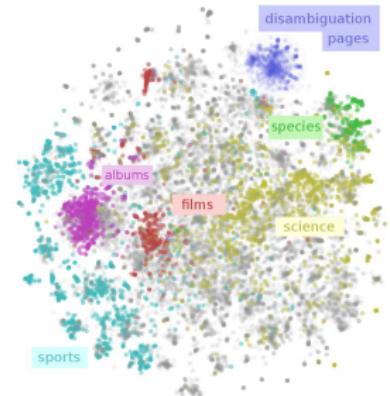
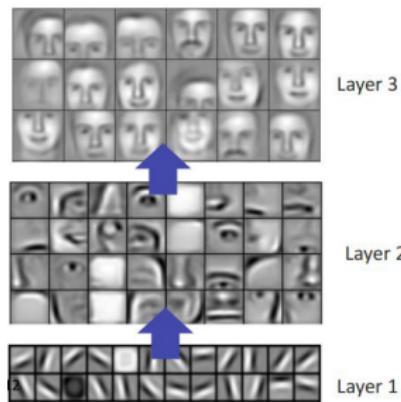
  Applications

  Open Problems

  CNN using Tensorflow

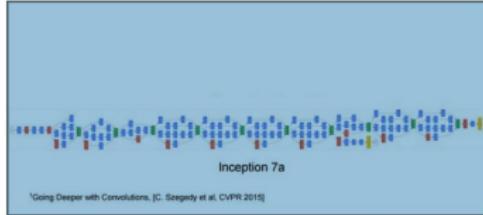
# Why Deep Learning?

- ▶ Hand-Engineered Features vs. Learned features



## Why Now?

- ▶ Large Datasets
- ▶ GPU Hardware Advances + Price Decreases
- ▶ Improved Techniques



## The Big Players



YAHOO!

Google



IBM



NVIDIA®

Baidu 百度

## Startups



vicarious



enlitic

clarifai

SKYMI<sup>N</sup>D

SIGNALSENSE

ersatz<sup>AI</sup>



Numen<sup>a</sup>ta



cortica<sup>TM</sup>



OpenAI

MetaMind



DEEPMIND

AlchemyAPI<sup>TM</sup>  
An IBM Company

wit.ai DNNresearch

# Outline

Introduction to Deep Learning

  Deep Learning Success

  Why Deep Learning and why now?

TensorFlow

**Introduction to TensorFlow**

    TensorFlow Example

    Linear Regression & Gradient Descent

Fundamentals of Deep Learning

  Deep Neural Network

  Artificial Neural Network

Computer Vision & Machine Learning

  Computer Vision in Machine Learning

Artificial Neural Networks (ANN)

  ANN for Computer Vision

  Neural Network for Image Classification

Image Classification using CNN

  Introduction to CNN

  Case Studies of CNN

  Applications

  Open Problems

  CNN using Tensorflow

# TensorFlow

- ▶ Open source software library for numerical computation using data flow graphs
- ▶ Originally developed by Google Brain Team to conduct machine learning and deep neural networks research
- ▶ General enough to be applicable in a wide variety of other domains as well
- ▶ TensorFlow provides an extensive suite of functions and classes that allow users to build various models from scratch

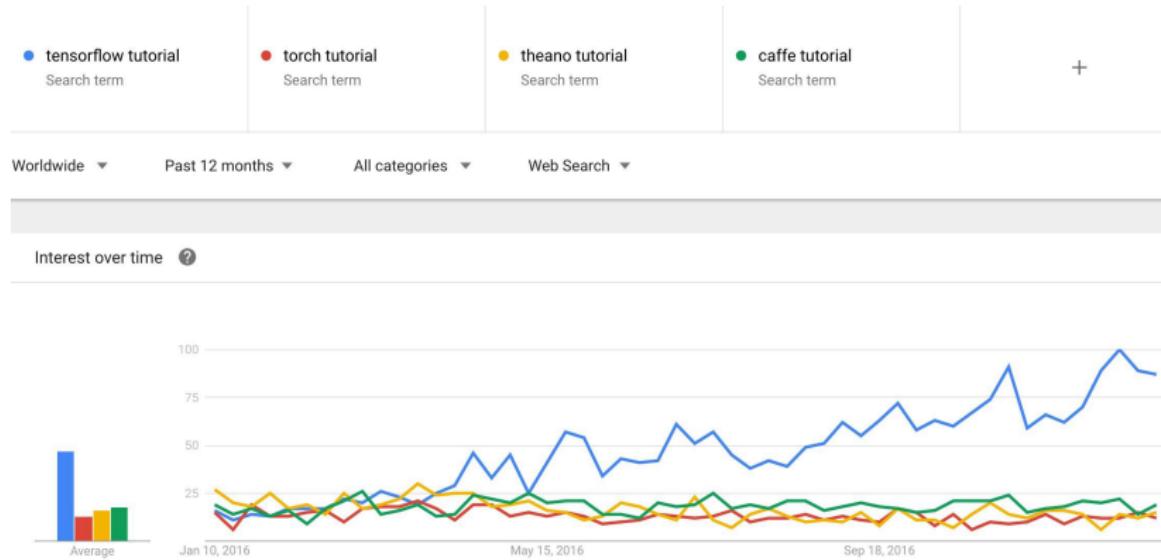


# Why TensorFlow?

- ▶ Python API
- ▶ Portability: deploy computation to one or more CPUs or GPUs in a desktop, server, or mobile device with a single API
- ▶ Auto-differentiation (no more taking derivatives by hand)
- ▶ Awesome projects already using TensorFlow
- ▶ Companies using Tensorflow
  - ▶ Google
  - ▶ OpenAI
  - ▶ DeepMind
  - ▶ Snapchat
  - ▶ Uber
  - ▶ Airbus
  - ▶ eBay
  - ▶ Dropbox
  - ▶ A bunch of startups

# Why TensorFlow?

- ▶ TF is not the only deep learning library



# Outline

Introduction to Deep Learning

  Deep Learning Success

  Why Deep Learning and why now?

**TensorFlow**

  Introduction to TensorFlow

**TensorFlow Example**

  Linear Regression & Gradient Descent

Fundamentals of Deep Learning

  Deep Neural Network

  Artificial Neural Network

Computer Vision & Machine Learning

  Computer Vision in Machine Learning

Artificial Neural Networks (ANN)

  ANN for Computer Vision

  Neural Network for Image Classification

Image Classification using CNN

  Introduction to CNN

  Case Studies of CNN

  Applications

  Open Problems

  CNN using Tensorflow

## TensorFlow Example

$$a = (b + c) * (c + 2) \quad (1)$$

- ▶ This function can be broken down into following operations:

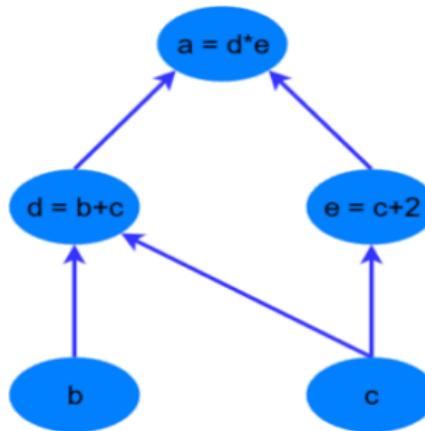
$$d = b + c \quad (2)$$

$$e = c + 2 \quad (3)$$

$$a = d * e \quad (4)$$

- ▶ Powerful idea in expressing the equation this way:
  - ▶ Two of the computations  $d$  and  $e$  can be performed in parallel
  - ▶ This can give us significant gains in computational times
  - ▶ These gains are a must for big data applications and deep learning

# Data Flow Graphs



- ▶ Data flows between different nodes in the graph are tensors which are multi-dimensional data arrays.

# What's a tensor?

- ▶ An n-dimensional array
  - ▶ 0-d tensor: scalar (number)
  - ▶ 1-d tensor: vector
  - ▶ 2-d tensor: matrix
  - ▶ and so on

## Simple TensorFlow example: Declaration Phase

$$a = (b + c) * (c + 2) \quad (5)$$

### Variable and Constant Declaration Code

```
import tensorflow as tf

# first, create a TensorFlow constant
const = tf.constant(2.0, name="const")

# create TensorFlow variables
b = tf.Variable(2.0, name='b')
c = tf.Variable(1.0, name='c')
```



## Simple TensorFlow Example Contd..

### TensorFlow Operations Code

```
# now create some operations
d = tf.add(b, c, name='d')
e = tf.add(c, const, name='e')
a = tf.multiply(d, e, name='a')
```

### Initialization Code

- ▶ Setup an object to initialise the variables and the graph structure.

```
# setup the variable initialisation
init_op = tf.global_variables_initializer()
```



## Simple TensorFlow Example Contd..

### Session Code

- ▶ To run the operations between the variables, we need to start a TensorFlow session. The TensorFlow session is an object where all operations are run.

```
# start the session
with tf.Session() as sess:
    # initialise the variables
    sess.run(init_op)
    # compute the output of the graph
    a_out = sess.run(a)
    print("Variable a is {}".format(a_out))
```



# TensorFlow Placeholder

## Placeholder

- ▶ To get the value of the array b during the `tf.Session()`, the variables are declared using the `tf.placeholder`.

```
# create TensorFlow variables  
b = tf.placeholder(tf.float32, name='b')
```

- ▶ The value may be passed as:  
`a_out = sess.run(a, feed_dict={b:5})`



## Visualize Graph with TensorBoard

```
import tensorflow as tf
a = tf.constant(2)
b = tf.constant(3)
x = tf.add(a, b)
with tf.Session() as sess:
    #Create the summary writer after graph definition and before
    running your session
    # add this line to use TensorBoard.
    writer = tf.summary.FileWriter('C:\\\\Graph')
    #Path where you want to keep your event files
    x_out= sess.run(x)
    print(x_out)
    writer.add_graph(sess.graph)
writer.close() # close the writer when you're done using it
```

## Run it

Go to terminal, run:

```
$ python [yourprogram].py  
$ tensorboard --logdir C:\Graph
```

Then open your browser and go to: <http://localhost:6006/>

**TensorBoard**

SCALARS

IMAGES

AUDIO

GRAPHS

DISTRIBUTIONS

HISTOGRAMS

EMBEDDINGS

Write a regex to create a tag group

 Split on underscores Data download links

Tooltip sorting method: default ▾



Go here

Smoothing



Horizontal Axis

STEP      RELATIVE      WALL

Runs

Write a regex to filter runs

  .

TOGGLE ALL RUNS

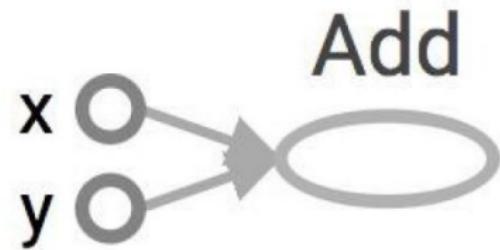
my\_graph

## More about Data Flow Graphs

```
import tensorflow as tf  
a = tf.add(3, 5)
```

- ▶ Why x, y?
- ▶ TF automatically names the nodes when you don't explicitly name them
  - ▶  $x=3$
  - ▶  $y=5$

Visualized by TensorBoard

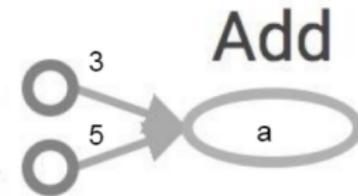


# More about Data Flow Graphs

Interpreted?

```
import tensorflow as tf  
a = tf.add(3, 5)
```

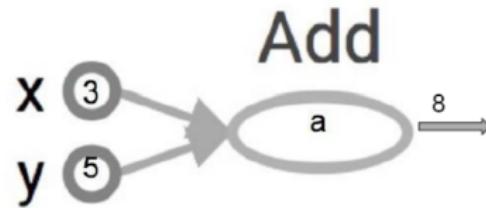
- ▶ Nodes: operators, variables, and constants
- ▶ Edges: tensors



## How to get the value of a?

- ▶ Create a session, assign it to variable sess so we can call it later
- ▶ Within the session, evaluate the graph to fetch the value of a

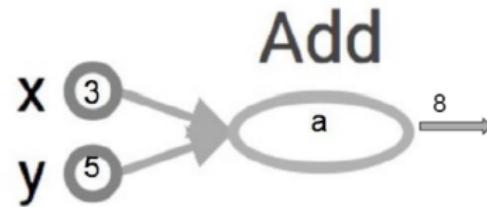
```
import tensorflow as tf  
a = tf.add(3, 5)  
sess = tf.Session()  
print sess.run(a)    >> 8  
sess.close()
```



## How to get the value of a?

- ▶ Create a session, assign it to variable sess so we can call it later
- ▶ Within the session, evaluate the graph to fetch the value of a

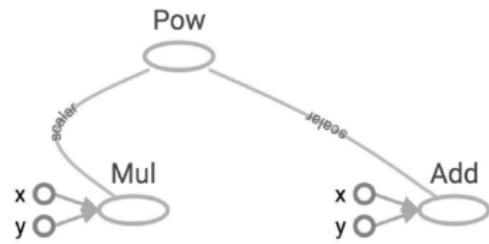
```
import tensorflow as tf  
a = tf.add(3, 5)  
sess = tf.Session()  
with tf.Session() as sess:  
    print sess.run(a)  
sess.close()
```



# More Graphs

```
import tensorflow as tf  
x = 2  
y = 3  
op1 = tf.add(x, y)  
op2 = tf.multiply(x, y)  
op3 = tf.pow(op2, op1)  
with tf.Session() as sess:  
    op3 = sess.run(op3)
```

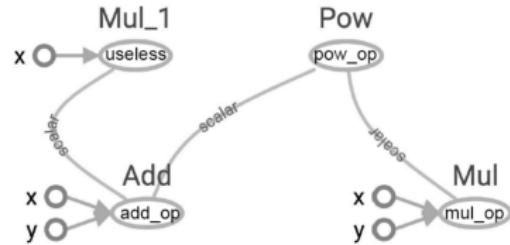
Visualized by TensorBoard



# Sub-Graphs

```
import tensorflow as tf
x = 2
y = 3
add_op = tf.add(x, y)
mul_op = tf.multiply(x, y)
useless = tf.multiply(x, add_op)
pow_op = tf.pow(add_op, mul_op)
with tf.Session() as sess:
    z = sess.run(pow_op)
```

Visualized by TensorBoard

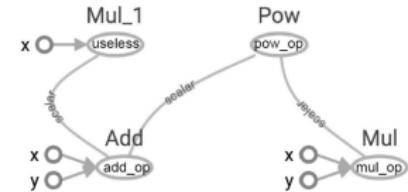


Because we only want the value of `pow_op` and `pow_op` doesn't depend on `useless`, session won't compute value of `useless`

→ save computation

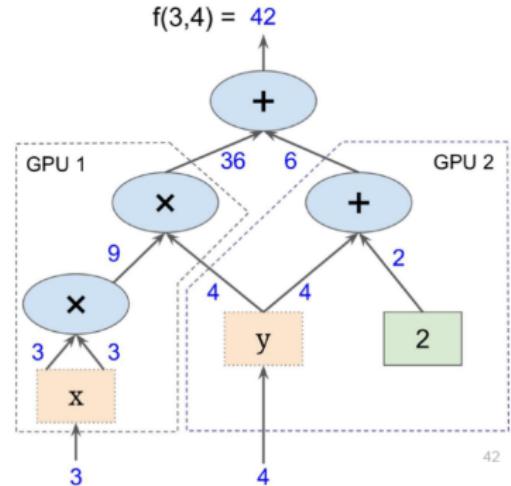
## Sub-Graphs

```
import tensorflow as tf
x = 2
y = 3
add_op = tf.add(x, y)
mul_op = tf.multiply(x, y)
useless = tf.multiply(x, add_op)
pow_op = tf.pow(add_op, mul_op)
with tf.Session() as sess:
    z, not_useless = sess.run([pow_op, useless])
```



# Sub-Graphs

- ▶ Possible to break graphs into several chunks and run them parallelly across multiple CPUs, GPUs, or devices



## More Constants

- ▶ `a = tf.constant([2, 2], name="a")` ← 1d Tensor
- ▶ `b = tf.constant([[0, 1], [2, 3]], name="b")` ← 2d Tensor
- ▶ Tensors filled with a specific value
  - ▶ `tf.zeros([2, 3], tf.int32)` ==> `[[0, 0, 0], [0, 0, 0]]`
  - ▶ Let `input_tensor` is `[0, 1], [2, 3], [4, 5]`  
`tf.zeros_like(input_tensor)` ==> `[[0, 0], [0, 0], [0, 0]]`
  - ▶ `tf.fill([2, 3], 8)` ==> `[[8, 8, 8], [8, 8, 8]]`



## More Constants

- ▶ Constants as sequences

- ▶ `tf.linspace(start, stop, num)`

`tf.linspace(10.0, 13.0, 4) ==> [10.0 11.0 12.0 13.0]`

- ▶ `tf.range(start, limit=None, delta=1, dtype=None, name='range')`

Let 'start' is 3, 'limit' is 18, 'delta' is 3

`tf.range(start, limit, delta) ==> [3, 6, 9, 12, 15]`

Let 'limit' is 5

`tf.range(limit) ==> [0, 1, 2, 3, 4]`

Tensor objects are not iterable

`for _ in tf.range(4): ==># TypeError`

# TensorFlow Data Types

| Data type     | Python type                | Description   |
|---------------|----------------------------|---|
| DT_FLOAT      | <code>tf.float32</code>    | 32 bits floating point.   |
| DT_DOUBLE     | <code>tf.float64</code>    | 64 bits floating point.   |
| DT_INT8       | <code>tf.int8</code>       | 8 bits signed integer.  |
| DT_INT16      | <code>tf.int16</code>      | 16 bits signed integer.   |
| DT_INT32      | <code>tf.int32</code>      | 32 bits signed integer.   |
| DT_INT64      | <code>tf.int64</code>      | 64 bits signed integer.   |
| DT_UINT8      | <code>tf.uint8</code>      | 8 bits unsigned integer.  |
| DT_UINT16     | <code>tf.uint16</code>     | 16 bits unsigned integer.   |
| DT_STRING     | <code>tf.string</code>     | Variable length byte arrays. Each element of a Tensor is a byte array.        |
| DT_BOOL       | <code>tf.bool</code>       | Boolean.  |
| DT_COMPLEX64  | <code>tf.complex64</code>  | Complex number made of two 32 bits floating points: real and imaginary parts. |
| DT_COMPLEX128 | <code>tf.complex128</code> | Complex number made of two 64 bits floating points: real and imaginary parts. |

# Outline

Introduction to Deep Learning

  Deep Learning Success

  Why Deep Learning and why now?

TensorFlow

  Introduction to TensorFlow

  TensorFlow Example

  Linear Regression & Gradient Descent

Fundamentals of Deep Learning

  Deep Neural Network

  Artificial Neural Network

Computer Vision & Machine Learning

  Computer Vision in Machine Learning

Artificial Neural Networks (ANN)

  ANN for Computer Vision

  Neural Network for Image Classification

Image Classification using CNN

  Introduction to CNN

  Case Studies of CNN

  Applications

  Open Problems

  CNN using Tensorflow

# Linear Regression & Gradient Descent

## ► Notation:

- ▶  $m$  : Number of training examples
- ▶  $x$  : Input variables (Features)
- ▶  $y$ : Output variables (Targets)
- ▶  $(x,y)$ : Training Example (Represents 1 row on the table)
- ▶  $(x^{(i)}, y^{(i)})$  :  $i^{th}$  training example (Represents's  $i^{th}$  row on the table)
- ▶  $n$  : Number of features (Dimensionality of the input)



## Representation of the Hypothesis (Function):

- ▶ In this case,  $y$  is represented as a linear combination of the inputs ( $x$ )
- ▶ Which leads to:

$$h(\theta) = \theta_0 + \theta_1 x_1 + \theta_2 x_2 \cdots + \theta_n x_n$$

where,  $\theta_i$  are the parameters (also called weights).

- ▶ For convenience and ease of representation:  $x_0 = 1$
- ▶ So that, the above equation becomes:

$$h(x) = \sum_{i=0}^n (\theta^T x)$$

- ▶ The objective now, is to 'learn' the parameters  $\theta$
- ▶ So that  $h(x)$  becomes as close to ' $y$ ' at least for the training set.

# Linear Regression & Gradient Descent

- ▶ Define a function that measures for each value of the theta's, how close the  $h(x^{(i)})$  are to the corresponding  $y^{(i)}$
- ▶ The 'cost function':

$$J(\theta) = \frac{1}{2} \sum_{i=1}^m (h_{(\theta)}(x^{(i)}) - y^{(i)})^2$$

- ▶ Choose the parameters so as to minimize  $J(\theta)$

## Linear Regression & Gradient Descent

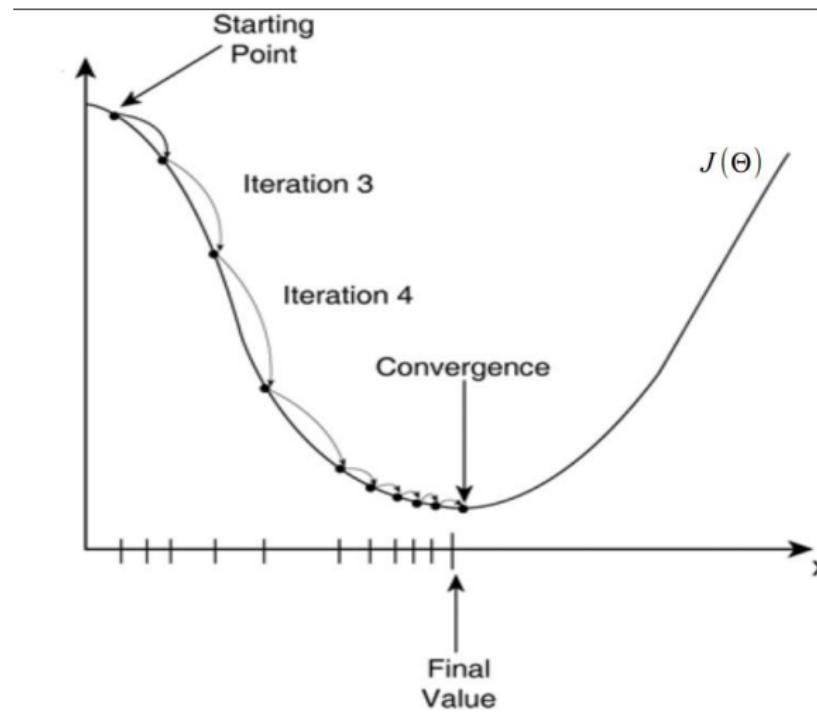
- ▶ To minimize, Gradient Descent algorithm may be applied
- ▶ Gradient Descent starts off with some initial  $\theta$ , and continually performs the following update:

$$\theta_j = \theta_j - \alpha \frac{\partial}{\partial \theta_j} J(\theta)$$

- ▶ This update is simultaneously performed for all values of  $j = 0, 1, \dots, n$
- ▶  $\alpha$  is called the learning rate



# Linear Regression & Gradient Descent



# TensorFlow Code for Linear Regression & Gradient Descent

## Import necessary libraries

```
import tensorflow as tf  
  
import numpy  
  
import matplotlib.pyplot as plt  
  
rng = numpy.random
```

# TensorFlow Code for Linear Regression & Gradient Descent

## Parameters

```
learning_rate = 0.01
```

```
training_epochs = 1000
```

```
display_step = 50
```

## Training Data

```
train_X = numpy.asarray([3.3,4.4,5.5,6.71,6.93,4.168,9.779,6.182,  
    7.59,2.167,7.042,10.791,5.313,7.997,5.654,9.27,3.1])
```

```
train_Y = numpy.asarray([1.7,2.76,2.09,3.19,1.694,1.573,3.366,  
    2.596,2.53,1.221,2.827,3.465,1.65,2.904,2.42,2.94,1.3])
```

```
n_samples = train_X.shape[0]
```

# TensorFlow Code for Linear Regression & Gradient Descent

## tf Graph Input

```
X = tf.placeholder("float")  
Y = tf.placeholder("float")
```

## Set model weights

```
W = tf.Variable(rng.randn(), name="weight")  
b = tf.Variable(rng.randn(), name="bias")
```



# TensorFlow Code for Linear Regression & Gradient Descent

Construct a linear model

```
pred = tf.add(tf.multiply(X, W), b)
```

Mean squared error

```
cost = tf.reduce_sum(tf.pow(pred-Y, 2))/(2*n_samples)
```

Gradient descent

```
optimizer =  
tf.train.GradientDescentOptimizer(learning_rate).minimize(cost)
```

Initialize the variables

```
init = tf.global_variables_initializer()
```

## Start training

```
with tf.Session() as sess:  
    sess.run(init) # Run the initializer  
    # Fit all training data  
    for epoch in range(training_epochs):  
        for (x, y) in zip(train_X, train_Y):  
            sess.run(optimizer, feed_dict={X: x, Y: y})  
    # Display logs per epoch step  
    if (epoch+1) % display_step == 0:  
        c = sess.run(cost, feed_dict={X: train_X, Y:train_Y})  
        print("Epoch:", '%04d' % (epoch+1), "cost=",  
              " {:.9f} ".format(c),"W=", sess.run(W), "b=",  
              sess.run(b))  
    print("Optimization Finished!")  
    training_cost = sess.run(cost, feed_dict=X: train_X, Y: train_Y)  
    print("Training cost=", training_cost, "W=", sess.run(W),  
          "b=", sess.run(b), '\n')
```

## Start training cont...

```
# Graphic display  
plt.plot(train_X, train_Y, 'ro', label='Original data')  
plt.plot(train_X, sess.run(W) * train_X + sess.run(b), label='Fitted  
line')  
plt.legend()  
plt.show()
```

## Testing example

```
test_X = numpy.asarray([6.83, 4.668, 8.9, 7.91, 5.7, 8.7, 3.1, 2.1])
test_Y = numpy.asarray([1.84, 2.273, 3.2, 2.831, 2.92, 3.24, 1.35,
1.03])
print("Testing... (Mean square loss Comparison)")
testing_cost = sess.run(
    tf.reduce_sum(tf.pow(pred - Y, 2)) / (2 * test_X.shape[0]),
    feed_dict={X: test_X, Y: test_Y} # same function as cost above
)
print("Testing cost=", testing_cost)
print("Absolute mean square loss difference:", abs(
    training_cost - testing_cost))
plt.plot(test_X, test_Y, 'bo', label='Testing data')
plt.plot(train_X, sess.run(W) * train_X + sess.run(b), label='Fitted
line')
plt.legend()
plt.show()
```

# Outline

Introduction to Deep Learning

  Deep Learning Success

  Why Deep Learning and why now?

TensorFlow

  Introduction to TensorFlow

  TensorFlow Example

  Linear Regression & Gradient Descent

Fundamentals of Deep Learning

  Deep Neural Network

  Artificial Neural Network

Computer Vision & Machine Learning

  Computer Vision in Machine Learning

Artificial Neural Networks (ANN)

  ANN for Computer Vision

  Neural Network for Image Classification

Image Classification using CNN

  Introduction to CNN

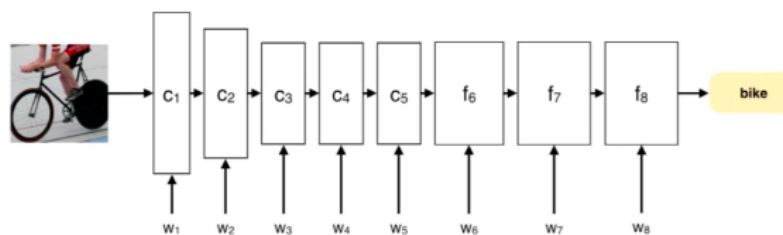
  Case Studies of CNN

  Applications

  Open Problems

  CNN using Tensorflow

# Deep Network: a magic box



# Outline

Introduction to Deep Learning

  Deep Learning Success

  Why Deep Learning and why now?

TensorFlow

  Introduction to TensorFlow

  TensorFlow Example

  Linear Regression & Gradient Descent

Fundamentals of Deep Learning

  Deep Neural Network

  Artificial Neural Network

Computer Vision & Machine Learning

  Computer Vision in Machine Learning

Artificial Neural Networks (ANN)

  ANN for Computer Vision

  Neural Network for Image Classification

Image Classification using CNN

  Introduction to CNN

  Case Studies of CNN

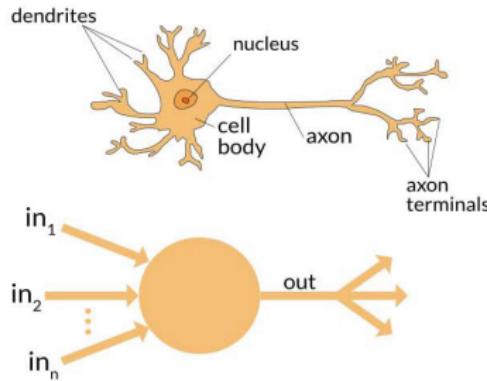
  Applications

  Open Problems

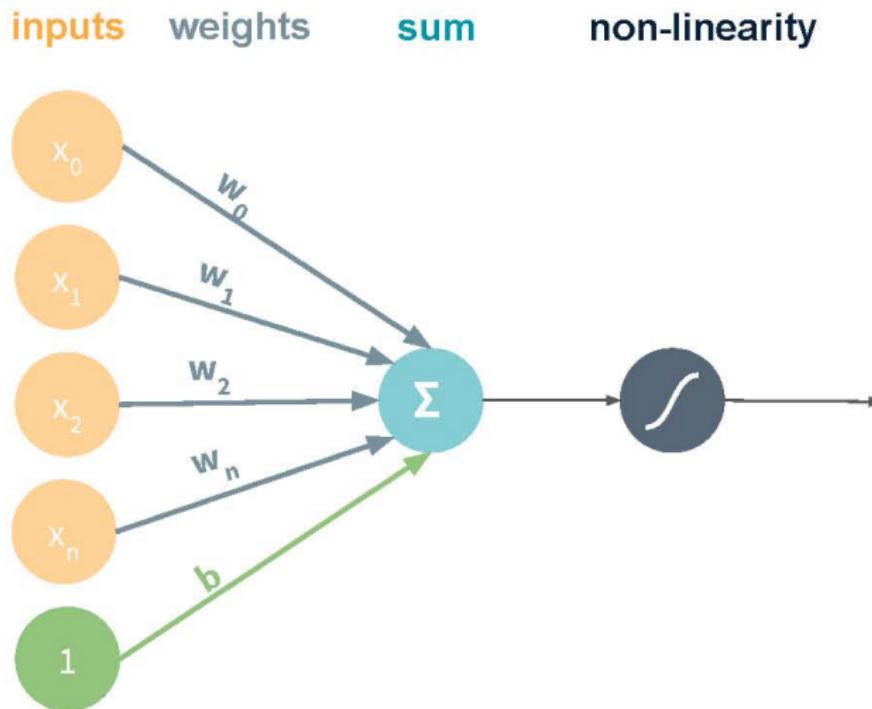
  CNN using Tensorflow

# The Perceptron

- ▶ Invented in 1954 by Frank Rosenblatt
- ▶ Inspired by neurobiology

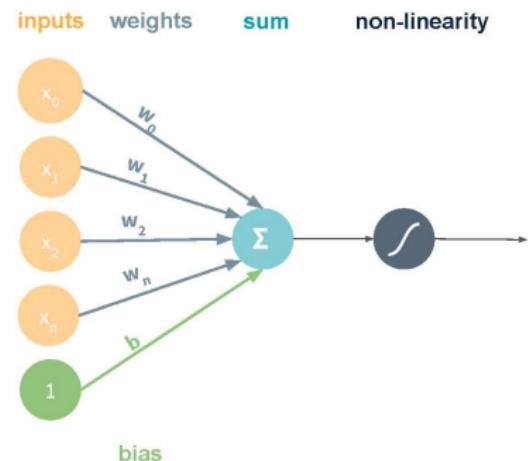


# The Perceptron



# Perceptron Forward Pass

$$\text{output} = g\left(\left(\sum_{i=0}^N x_i * w_i\right) + b\right)$$

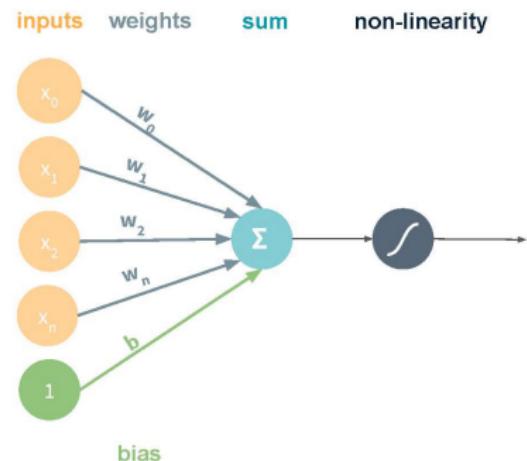


# Perceptron Forward Pass

$$\text{output} = g(XW + b)$$

$$X = x_0, x_1, \dots, x_n$$

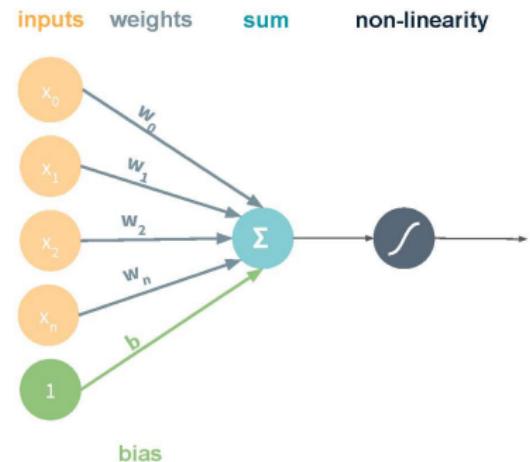
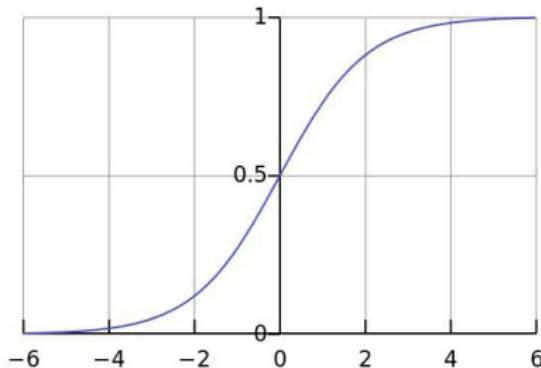
$$W = w_0, w_1, \dots, w_n$$



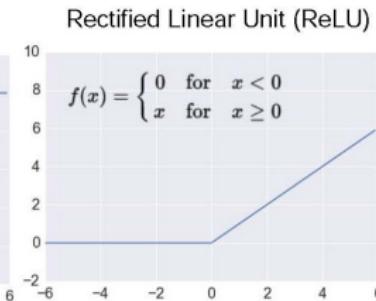
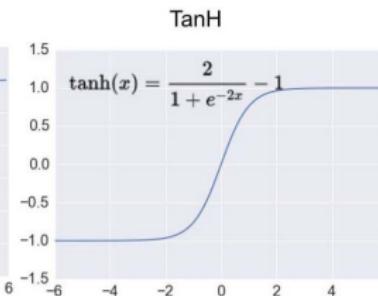
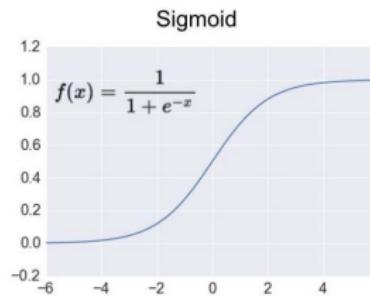
# Sigmoid Activation

$$\text{output} = g(XW + b)$$

$$g(a) = \sigma(a) = \frac{1}{1 + e^{-a}}$$

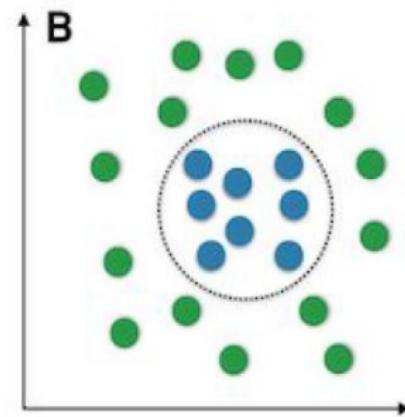
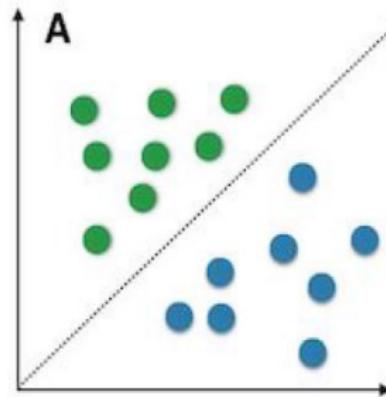


# Common Activation Functions



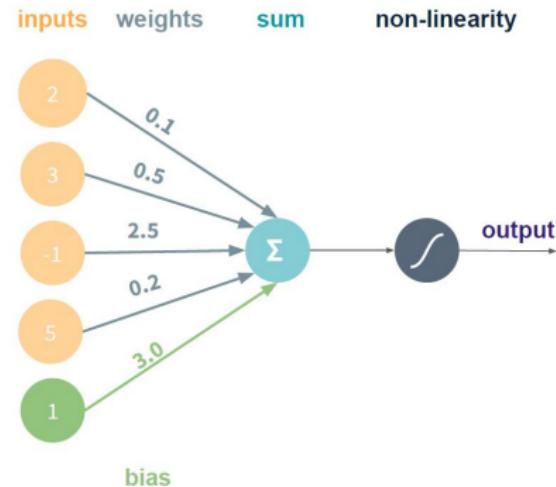
## Importance of Activation Functions

- ▶ Activation functions add non-linearity to our networks function
- ▶ Most real-world problems + data are non-linear



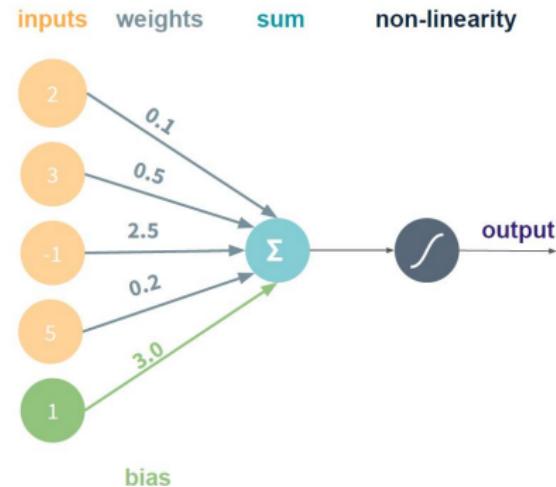
# Perceptron Forward Pass

$$\text{output} = g(XW + b)$$



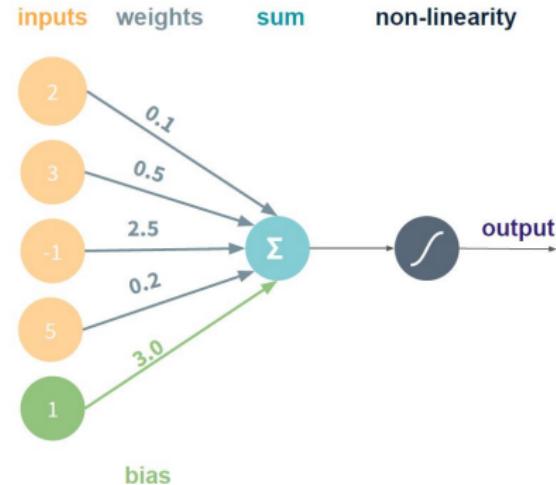
# Perceptron Forward Pass

$$\text{output} = g(\text{(sum of weighted inputs + bias)})$$
$$\begin{aligned} & (2 * 0.1) + \\ & (3 * 0.5) + \\ & (-1 * 2.5) + \\ & (5 * 0.2) + \\ & (1 * 3.0) \end{aligned}$$



# Perceptron Forward Pass

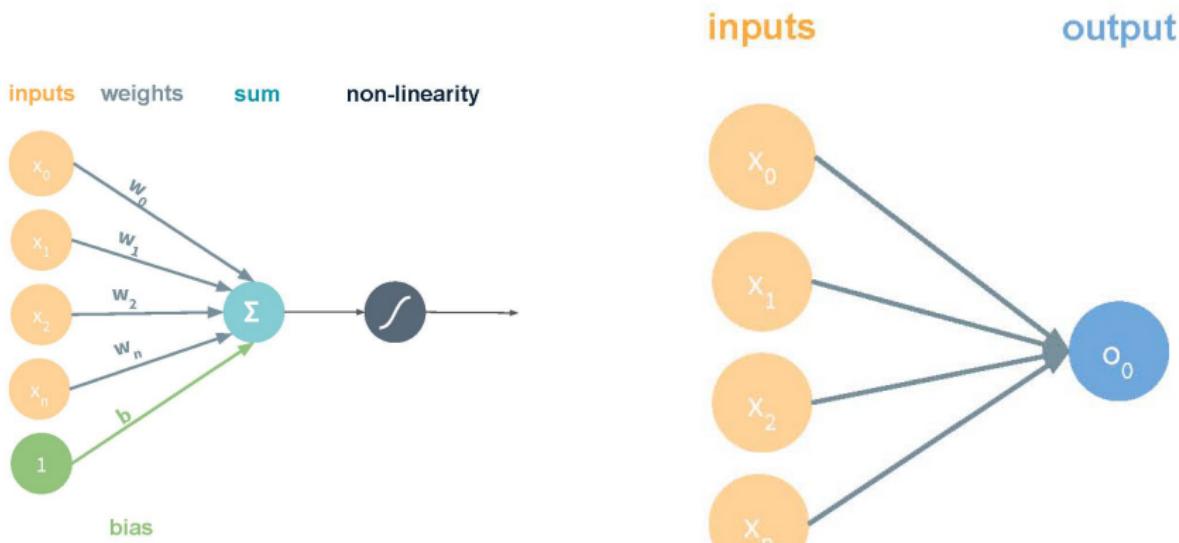
$$\begin{aligned} \text{output} &= g(3.2) = \sigma(3.2) \\ &= \frac{1}{(1 + e^{-3.2})} = 0.96 \end{aligned}$$



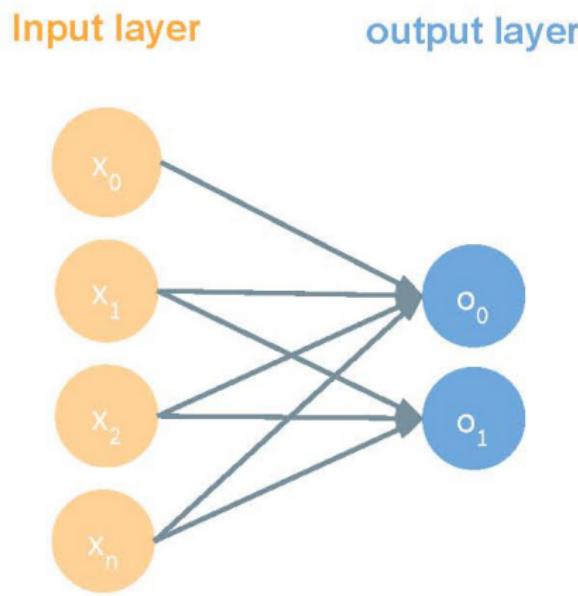
# How do we build neural networks with perceptrons?



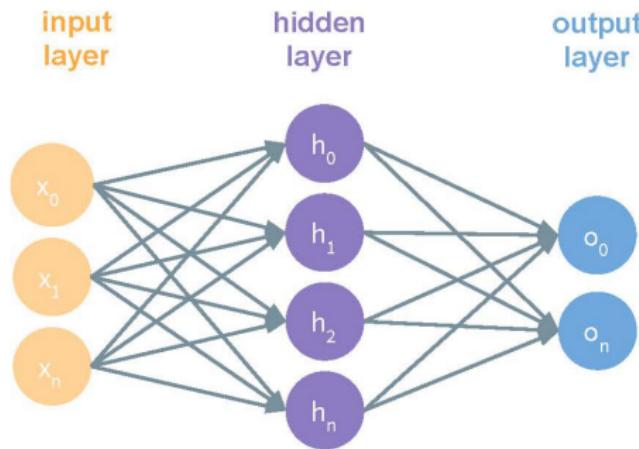
# Perceptron Diagram Simplified



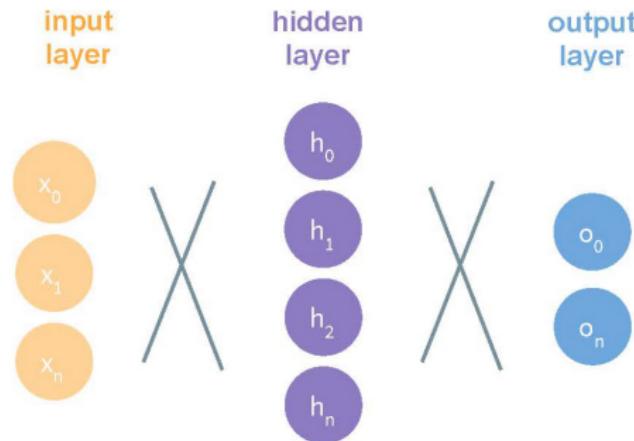
# Multi-Output Perceptron



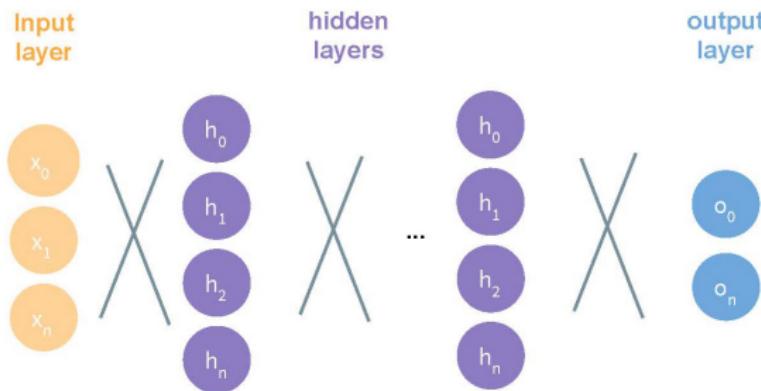
# Multi-Layer Perceptron (MLP)



# Multi-Layer Perceptron (MLP)



# Deep Neural Network



# Applying Neural Networks

## Example Problem: Will my Flight be Delayed?



## Example Problem: Will my Flight be Delayed?

**Temperature: -20 F**

**Wind Speed: 45 mph**



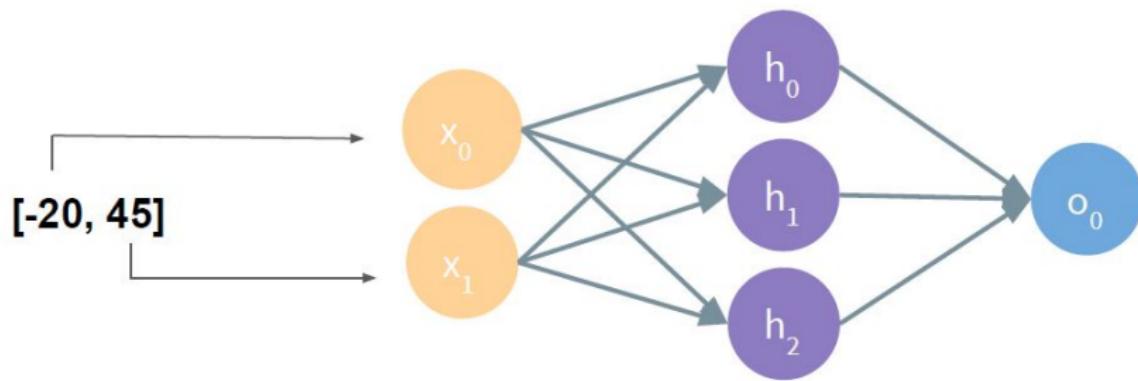
## Example Problem: Will my Flight be Delayed?

**[-20, 45]**

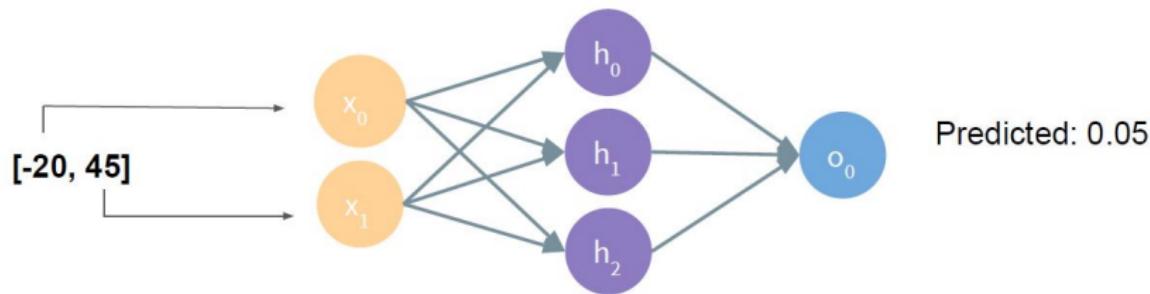


| DEPARTURE TIME | DESTINATION | STATUS  |
|----------------|-------------|---------|
| 0805 0815      | MILANO      | DELAYED |
| 0845 0855      | PARIS       | DELAYED |
| 0915 0915      | NEW YORK    | DELAYED |
| 0920 0920      | FRANKFURT   | DELAYED |
| 0925 0935      | LONDON      | DELAYED |
|                |             | DELAYED |

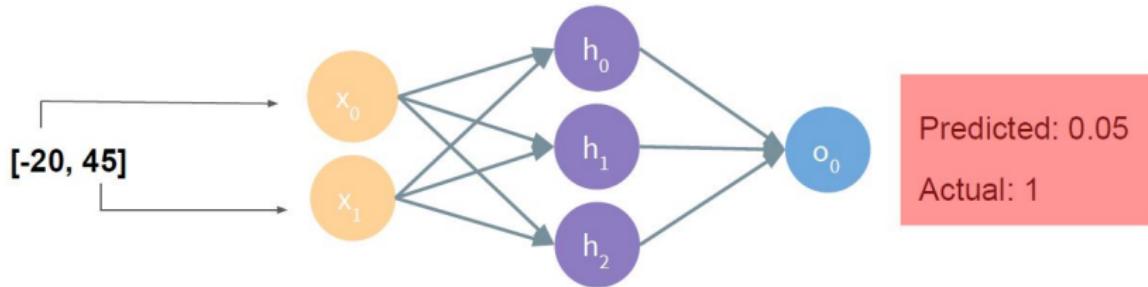
## Example Problem: Will my Flight be Delayed?



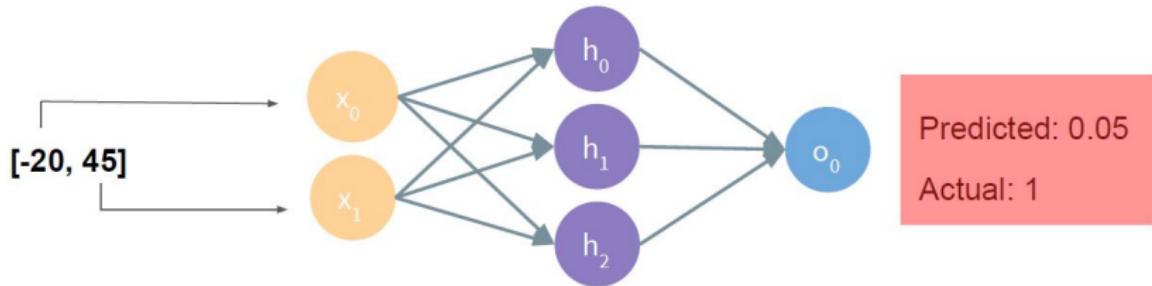
## Example Problem: Will my Flight be Delayed?



# Example Problem: Will my Flight be Delayed?



# Quantifying Loss

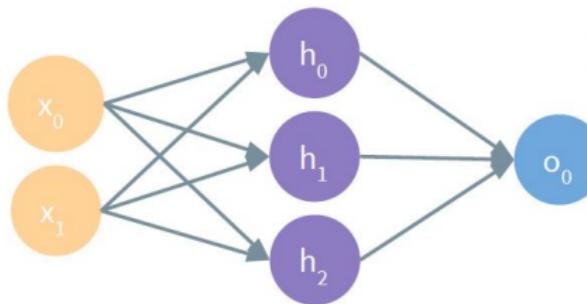


$$\text{loss}(f(x^{(i)}; \theta), y^{(i)})$$

# Total Loss

## Input

```
[  
[-20, 45],  
[80, 0],  
[4, 15],  
[45, 60],  
]
```



## Predicted

```
[  
0.05  
0.02  
0.96  
0.35  
]
```

## Actual

```
[  
1  
0  
1  
1  
]
```

$$\text{Total loss} := J(\theta) = \frac{1}{N} \sum_i^N \text{loss}(f(x^{(i)}; \theta), y^{(i)})$$

# Training Neural Networks

# Training Neural Networks: Objective

$$\arg_{\theta} \min \frac{1}{N} \sum_i^N loss(f(x^{(i)}; \theta), y^{(i)})$$

## Training Neural Networks: Objective

$$\arg_{\theta} \min \frac{1}{N} \sum_i^N \text{loss}(f(x^{(i)}; \theta), y^{(i)})$$



$J(\theta)$



**loss function**

## Training Neural Networks: Objective

$$\arg_{\theta} \min \frac{1}{N} \sum_i^N \text{loss}(f(x^{(i)}; \theta), y^{(i)})$$

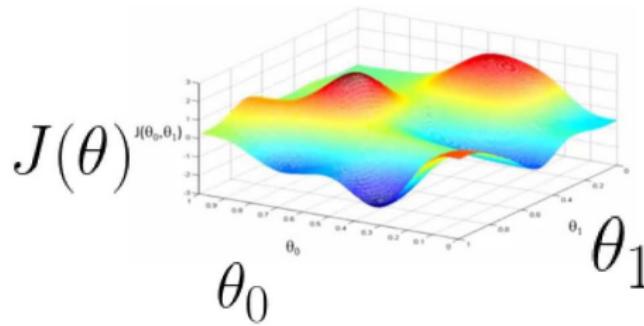


$$J(\theta) \leftarrow$$

$$\theta = W_1, W_2 \dots W_n$$

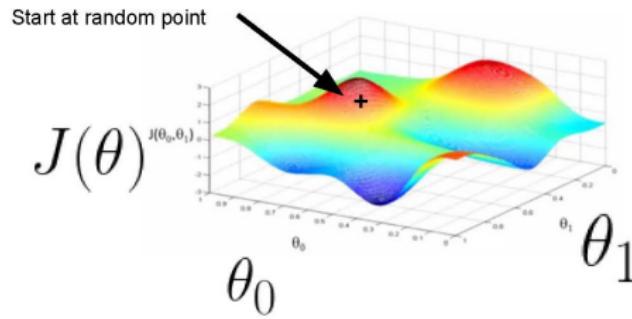
# Training Neural Networks: Objective

- ▶ Loss is a function of the models parameters



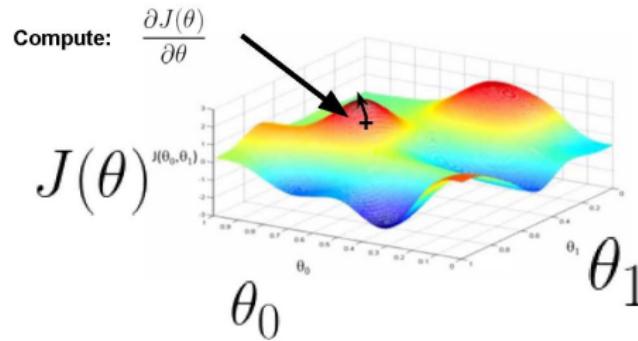
# Training Neural Networks: Objective

- ▶ How to minimize loss?



# Training Neural Networks: Objective

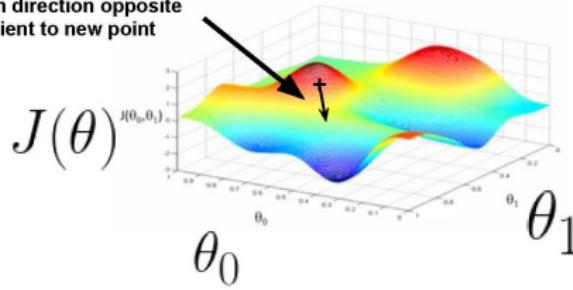
- ▶ How to minimize loss?



# Training Neural Networks: Objective

- ▶ How to minimize loss?

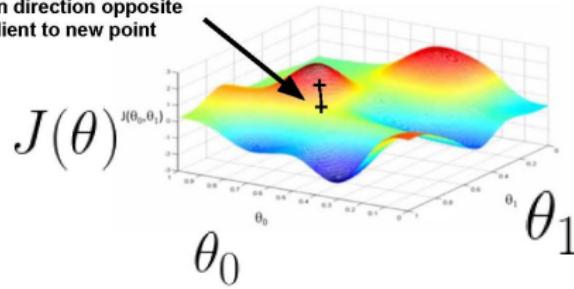
Move in direction opposite  
of gradient to new point



# Training Neural Networks: Objective

- ▶ How to minimize loss?

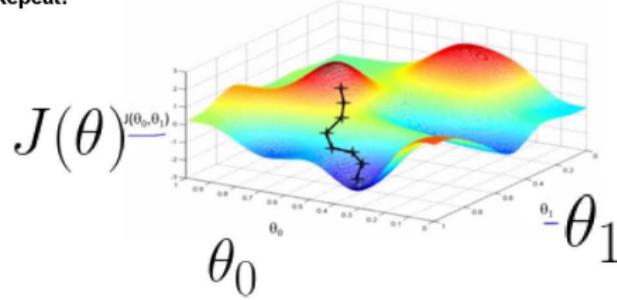
Move in direction opposite  
of gradient to new point



# Training Neural Networks: Objective

- ▶ How to minimize loss?

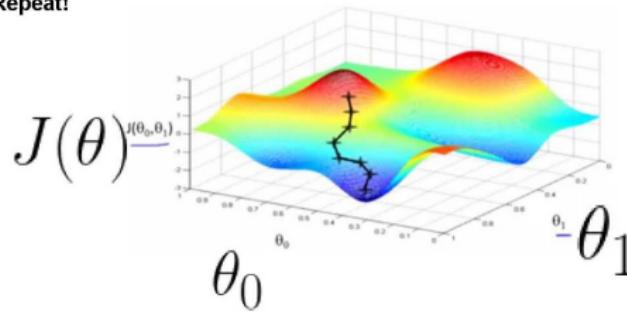
Repeat!



# Training Neural Networks: Objective

- ▶ This is called Stochastic Gradient Descent (SGD)

Repeat!



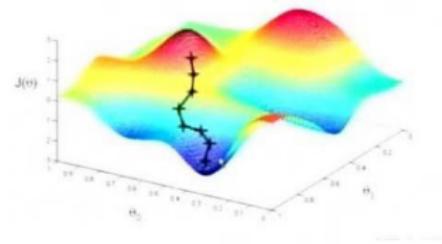
# Stochastic Gradient Descent (SGD)

- ▶ Initialize  $\theta$  randomly
- ▶ For N Epochs
  - ▶ For each training example  $(x, y)$ :
    - ▶ Compute Loss Gradient:

$$\frac{\partial J(\theta)}{\partial \theta}$$

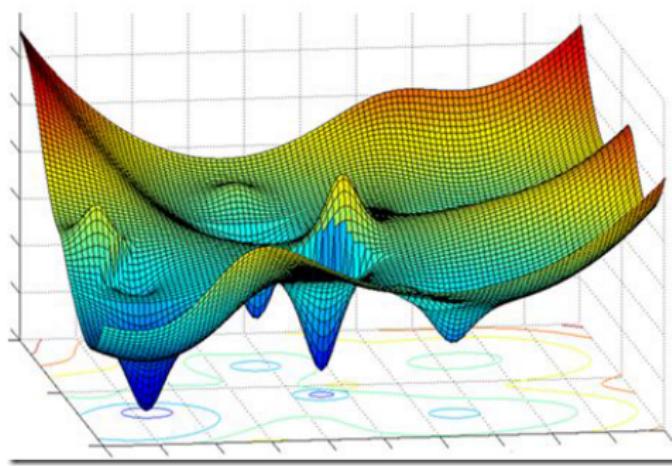
- ▶ Update  $\theta$  with update rule:

$$\theta := \theta - \eta \frac{\partial J(\theta)}{\partial \theta}$$



# Training Neural Networks In Practice

## Loss function can be difficult to optimize



# Loss function can be difficult to optimize

How to Choose Learning Rate?

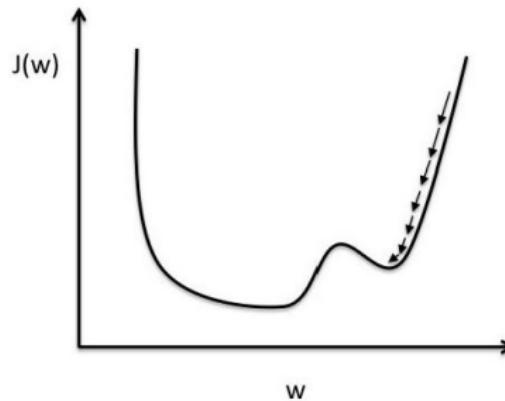
**Update Rule:**

$$\theta := \theta - \eta \frac{\partial J(\theta)}{\partial \theta}$$



# Learning Rate & Optimization

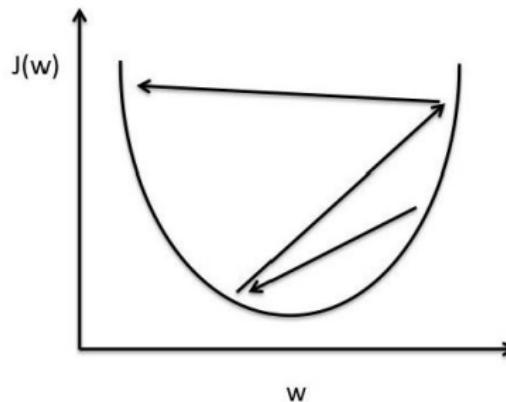
- ▶ Small Learning Rate



**Small learning rate: Many iterations until convergence and trapping in local minima.**

# Learning Rate & Optimization

- ▶ Large Learning Rate



Large learning rate: Overshooting.

## How to deal with this?

- ▶ Try lots of different learning rates to see what is just right

## How to deal with this?

- ▶ Try lots of different learning rates to see what is just right
- ▶ Do something smarter



## How to deal with this?

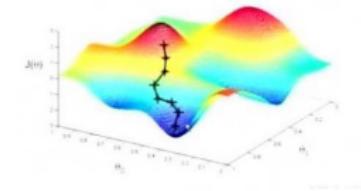
- ▶ Try lots of different learning rates to see what is just right
- ▶ **Do something smarter : Adaptive Learning Rate**
  - ▶ ADAM
  - ▶ Momentum
  - ▶ NAG
  - ▶ Adagrad
  - ▶ Adadelta
  - ▶ RMSProp
- ▶ For details: check out  
<http://sebastianruder.com/optimizing-gradient-descent/>

## Training Neural Networks in Practice 2: MiniBatches

# Why is it Stochastic Gradient Descent?

- Initialize  $\theta$  randomly
- For N Epochs
  - For each training example  $(x, y)$ :
    - Compute Loss Gradient:  $\frac{\partial J(\theta)}{\partial \theta}$
    - Update  $\theta$  with update rule:

Only an estimate of  
true gradient!



$$\theta := \theta - \eta \frac{\partial J(\theta)}{\partial \theta}$$

# Minibatches Reduce Gradient Variance

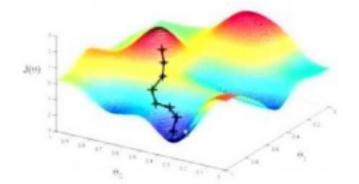
- Initialize  $\theta$  randomly
- For N Epochs

- For each training batch  $\{(x_0, y_0), \dots, (x_B, y_B)\}$ :

- Compute Loss Gradient: 
$$\frac{\partial J(\theta)}{\partial \theta} = \frac{1}{B} \sum_i^B \frac{\partial J_i(\theta)}{\partial \theta}$$
- Update  $\theta$  with update rule:

$$\theta := \theta - \eta \frac{\partial J(\theta)}{\partial \theta}$$

More accurate estimate!



## Training Neural Networks In Practice 3: Fighting Overfitting

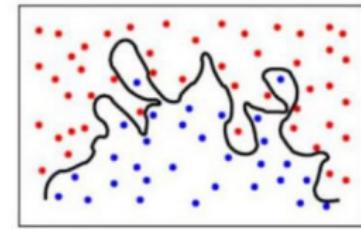
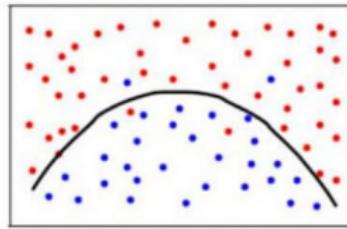
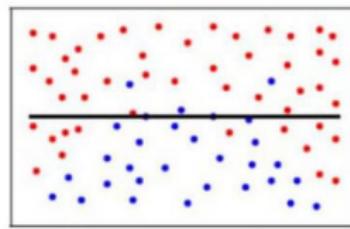


# The Problem of Overfitting

Underfitting

↔

Overfitting



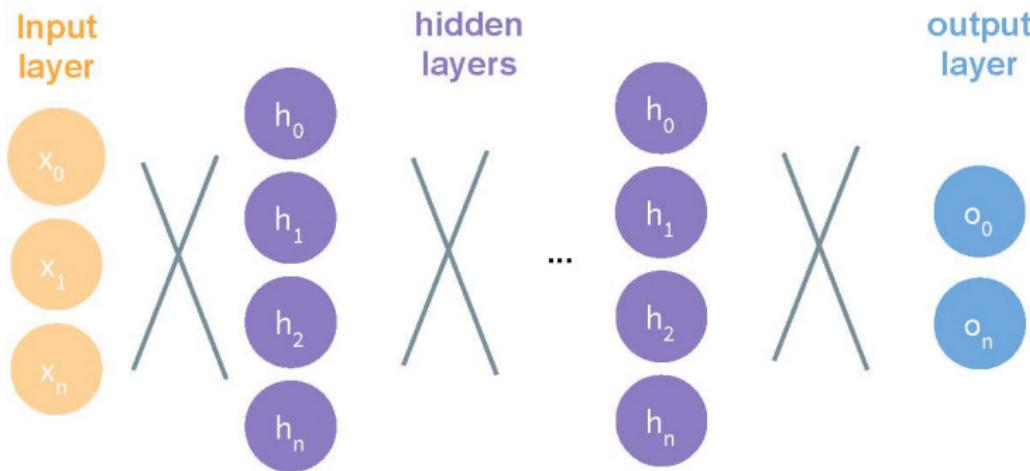
## Popular Regularization Techniques

- ▶ Dropout
- ▶ Early Stopping



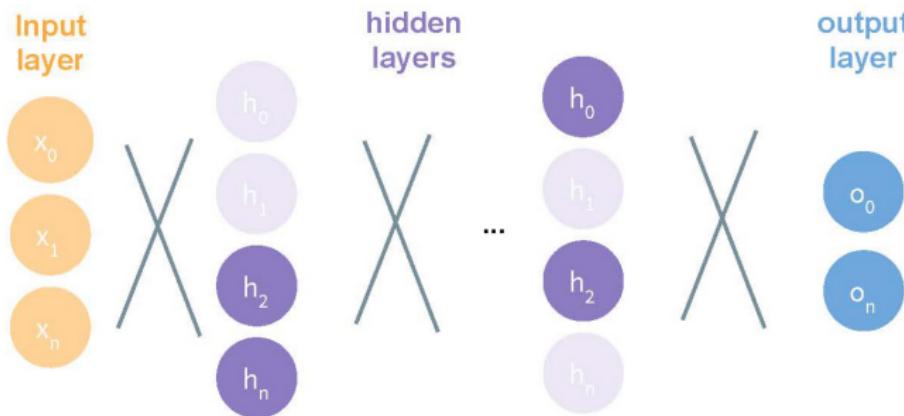
# Regularization I: Dropout

- ▶ During training, randomly set some activations to 0



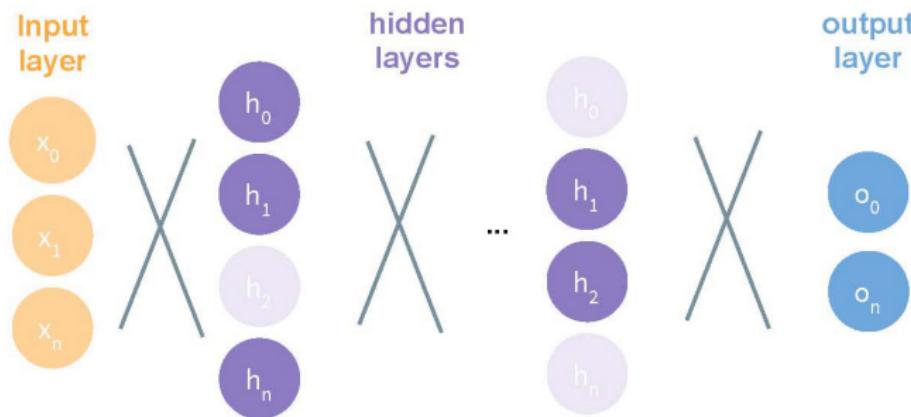
# Regularization I: Dropout

- ▶ During training, randomly set some activations to 0
  - ▶ Typically drop 50% of activations in layer
  - ▶ Forces network to not rely on any 1 node



# Regularization I: Dropout

- ▶ During training, randomly set some activations to 0
  - ▶ Typically drop 50% of activations in layer
  - ▶ Forces network to not rely on any 1 node



## Regularization II: Early Stopping

- ▶ Dont give the network time to overfit
- ▶ ...
- ▶ Epoch 15: Train: 85% Validation: 80%
- ▶ Epoch 16: Train: 87% Validation: 82%
- ▶ Epoch 17: Train: 90% Validation: 85%
- ▶ Epoch 18: Train: 95% **Validation: 83%**
- ▶ Epoch 19: Train: 97% **Validation: 78%**
- ▶ Epoch 20: Train: 98% **Validation: 75%**

## Regularization II: Early Stopping

- ▶ Dont give the network time to overfit
- ▶ ...
- ▶ Epoch 15: Train: 85% Validation: 80%
- ▶ Epoch 16: Train: 87% Validation: 82%
- ▶ Epoch 17: **Train: 90% Validation: 85%** ← Stop Here
- ▶ Epoch 18: Train: 95% **Validation: 83%**
- ▶ Epoch 19: Train: 97% **Validation: 78%**
- ▶ Epoch 20: Train: 98% **Validation: 75%**



# Outline

Introduction to Deep Learning

  Deep Learning Success

  Why Deep Learning and why now?

TensorFlow

  Introduction to TensorFlow

  TensorFlow Example

  Linear Regression & Gradient Descent

Fundamentals of Deep Learning

  Deep Neural Network

  Artificial Neural Network

Computer Vision & Machine Learning

  Computer Vision in Machine Learning

Artificial Neural Networks (ANN)

  ANN for Computer Vision

  Neural Network for Image Classification

Image Classification using CNN

  Introduction to CNN

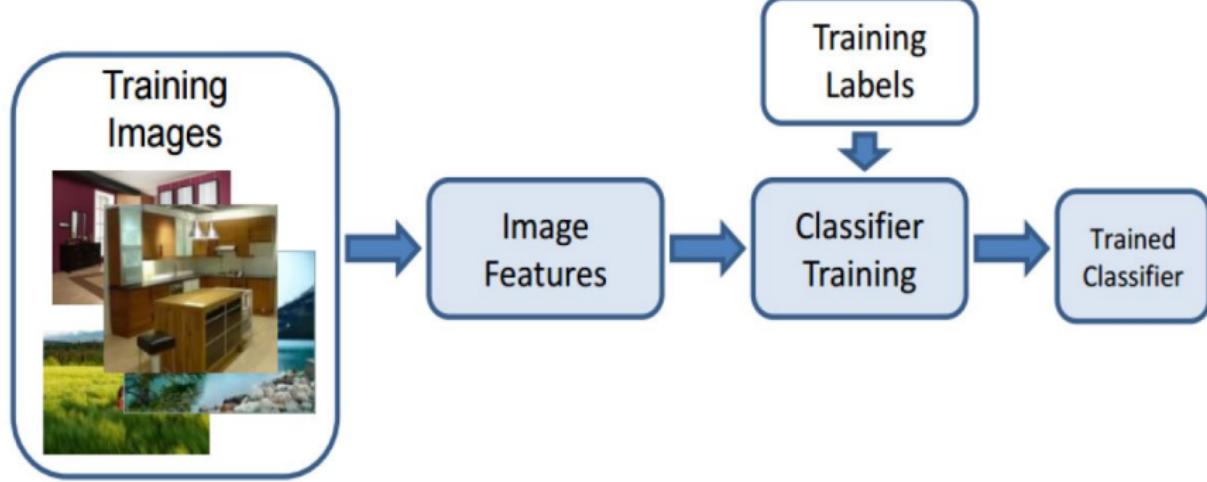
  Case Studies of CNN

  Applications

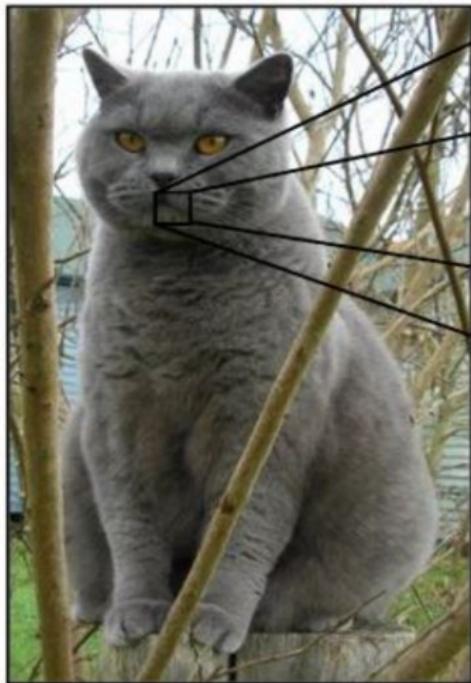
  Open Problems

  CNN using Tensorflow

# Computer Vision is Machine Learning



# Images are Numbers



|    |    |    |    |    |    |    |    |    |    |    |    |    |    |    |    |    |    |    |    |
|----|----|----|----|----|----|----|----|----|----|----|----|----|----|----|----|----|----|----|----|
| 08 | 02 | 22 | 97 | 38 | 15 | 00 | 40 | 00 | 75 | 04 | 05 | 07 | 78 | 52 | 12 | 50 | 79 | 01 | 50 |
| 49 | 49 | 99 | 40 | 17 | 81 | 18 | 57 | 60 | 87 | 17 | 40 | 98 | 43 | 65 | 18 | 01 | 56 | 62 | 00 |
| 81 | 49 | 31 | 73 | 55 | 79 | 14 | 29 | 93 | 71 | 40 | 67 | 21 | 85 | 30 | 03 | 49 | 13 | 36 | 65 |
| 52 | 70 | 95 | 23 | 04 | 60 | 11 | 42 | 62 | 11 | 05 | 56 | 01 | 32 | 56 | 71 | 37 | 02 | 36 | 91 |
| 22 | 31 | 16 | 71 | 51 | 62 | 03 | 59 | 91 | 92 | 36 | 54 | 22 | 40 | 40 | 28 | 66 | 33 | 13 | 80 |
| 24 | 47 | 34 | 00 | 99 | 03 | 45 | 02 | 44 | 75 | 33 | 53 | 78 | 36 | 84 | 20 | 35 | 17 | 12 | 50 |
| 12 | 95 | 81 | 28 | 64 | 23 | 67 | 10 | 26 | 38 | 40 | 67 | 59 | 54 | 70 | 66 | 18 | 38 | 64 | 70 |
| 67 | 26 | 20 | 68 | 02 | 62 | 12 | 20 | 95 | 63 | 94 | 39 | 63 | 08 | 40 | 91 | 66 | 49 | 94 | 21 |
| 24 | 55 | 58 | 05 | 66 | 73 | 99 | 26 | 97 | 17 | 78 | 78 | 96 | 83 | 14 | 88 | 34 | 69 | 63 | 72 |
| 21 | 36 | 23 | 09 | 75 | 00 | 76 | 44 | 20 | 45 | 35 | 14 | 00 | 61 | 33 | 97 | 34 | 31 | 33 | 95 |
| 78 | 17 | 53 | 28 | 22 | 75 | 31 | 47 | 15 | 94 | 03 | 80 | 04 | 62 | 16 | 14 | 09 | 53 | 56 | 92 |
| 16 | 39 | 05 | 42 | 96 | 35 | 31 | 47 | 55 | 58 | 88 | 24 | 00 | 17 | 54 | 24 | 36 | 29 | 85 | 57 |
| 86 | 56 | 00 | 48 | 35 | 71 | 89 | 07 | 05 | 44 | 44 | 37 | 44 | 60 | 21 | 58 | 51 | 54 | 17 | 58 |
| 19 | 80 | 81 | 68 | 05 | 94 | 47 | 69 | 28 | 73 | 92 | 13 | 86 | 52 | 17 | 77 | 04 | 89 | 55 | 40 |
| 04 | 52 | 08 | 83 | 97 | 35 | 99 | 16 | 07 | 97 | 57 | 32 | 16 | 26 | 26 | 79 | 33 | 27 | 98 | 66 |
| 04 | 46 | 48 | 87 | 57 | 62 | 20 | 72 | 03 | 46 | 33 | 67 | 46 | 55 | 12 | 32 | 63 | 93 | 53 | 69 |
| 04 | 42 | 16 | 73 | 21 | 93 | 39 | 11 | 24 | 94 | 72 | 18 | 08 | 46 | 29 | 32 | 40 | 62 | 76 | 36 |
| 20 | 69 | 36 | 41 | 72 | 30 | 23 | 88 | 34 | 78 | 89 | 69 | 82 | 67 | 59 | 85 | 74 | 04 | 36 | 16 |
| 20 | 73 | 35 | 29 | 78 | 31 | 90 | 01 | 74 | 31 | 49 | 71 | 48 | 64 | 11 | 16 | 23 | 57 | 05 | 54 |
| 01 | 70 | 54 | 71 | 83 | 51 | 54 | 69 | 16 | 92 | 33 | 48 | 61 | 43 | 52 | 01 | 89 | 13 | 42 | 48 |

What the computer sees

image classification

82% cat  
15% dog  
2% hat  
1% mug

# Computer Vision is Hard

Viewpoint variation



Scale variation



Deformation



Occlusion



Illumination conditions



Background clutter



Intra-class variation



# Famous Computer Vision Datasets

|   |   |   |   |   |   |   |   |   |   |   |   |   |   |   |   |   |   |   |   |   |
|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|
| 9 | 3 | 1 | 9 | 7 | 2 | 4 | 5 | / | 0 | 3 | 2 | 4 | 3 | 7 | 5 | 9 | 0 | 3 | 4 | 5 |
| 3 | 0 | 2 | 4 | 2 | 9 | 4 | 8 | 3 | 0 | 1 | 3 | 5 | 3 | 8 | 7 | 4 | 6 | 1 | 1 | 9 |
| 2 | 8 | 3 | 2 | 3 | 8 | 2 | 4 | 9 | 8 | 9 | 1 | 3 | 9 | 1 | 5 | 0 | 4 | 5 | 9 | 6 |
| 6 | 3 | 6 | 9 | 0 | 3 | 6 | 0 | 3 | 0 | 1 | 3 | 9 | 4 | 5 | 4 | 6 | 5 | 7 | 4 | 1 |
| 3 | 3 | 8 | 0 | 7 | 0 | 5 | 6 | 9 | 8 | 3 | 4 | 1 | 2 | 0 | 9 | 8 | 7 | 4 | 1 | 9 |
| 1 | 1 | 9 | 6 | 3 | 0 | 4 | 3 | 7 | 9 | 5 | 0 | 5 | 4 | 6 | 5 | 7 | 4 | 1 | 9 | 8 |
| 9 | 5 | 0 | 0 | 5 | 1 | 1 | 7 | 9 | 4 | 7 | 7 | 0 | 6 | 5 | 8 | 9 | 4 | 3 | 2 | 1 |
| 0 | 2 | 0 | 1 | 6 | 7 | 0 | 9 | 5 | 6 | 3 | 8 | 0 | 5 | 4 | 6 | 5 | 7 | 4 | 1 | 9 |
| 9 | 4 | 3 | 2 | 1 | 6 | 0 | 2 | 0 | 0 | 7 | 8 | 4 | 6 | 5 | 0 | 3 | 4 | 5 | 9 | 6 |
| 5 | 1 | 6 | 1 | 6 | 6 | 2 | 9 | 7 | 4 | 6 | 0 | 7 | 6 | 5 | 8 | 9 | 8 | 7 | 4 | 1 |
| 7 | 1 | 7 | 5 | 9 | 2 | 3 | 7 | 9 | 6 | 1 | 3 | 0 | 4 | 5 | 8 | 0 | 4 | 0 | 5 | 6 |
| 1 | 6 | 7 | 9 | 4 | 4 | 1 | 4 | 9 | 1 | 3 | 1 | 2 | 3 | 4 | 8 | 1 | 5 | 5 | 6 | 7 |
| 0 | 1 | 6 | 1 | 6 | 1 | 6 | 7 | 5 | 5 | 0 | 6 | 8 | 7 | 1 | 3 | 8 | 7 | 4 | 5 | 6 |
| 6 | 4 | 8 | 7 | 3 | 1 | 3 | 0 | 9 | 5 | 0 | 3 | 8 | 6 | 2 | 4 | 9 | 3 | 5 | 0 | 3 |
| 7 | 9 | 3 | 1 | 3 | 0 | 9 | 7 | 9 | 5 | 0 | 3 | 8 | 6 | 2 | 4 | 9 | 3 | 5 | 0 | 3 |

## MNIST: handwritten digits



## CIFAR-10(0): tiny images

german housing airport weight  
teacher computer dead headquarters  
register gallery court key structure  
king fireplace church press  
cousin road paper cut  
hotel room tree file side site door  
spider plant wall means for hill car  
plant house school railcar  
house school film stick  
longer table top man car study  
spring range lamp net menu  
people shop net half glass  
kitchen engine bin memorabilia ball  
diaper stone child case student  
apple gift rule chair rule  
bag home room office club  
radio support level line street golf  
beach library stage video food building  
tool material player machine security call  
football hospital scale equipment cell phone mountain ring telephone  
short circuit bags scale gas meter telephone  
gas meter telephone

## ImageNet: WordNet hierarchy



## Places: natural scenes

# Outline

Introduction to Deep Learning

  Deep Learning Success

  Why Deep Learning and why now?

TensorFlow

  Introduction to TensorFlow

  TensorFlow Example

  Linear Regression & Gradient Descent

Fundamentals of Deep Learning

  Deep Neural Network

  Artificial Neural Network

Computer Vision & Machine Learning

  Computer Vision in Machine Learning

Artificial Neural Networks (ANN)

  ANN for Computer Vision

  Neural Network for Image Classification

Image Classification using CNN

  Introduction to CNN

  Case Studies of CNN

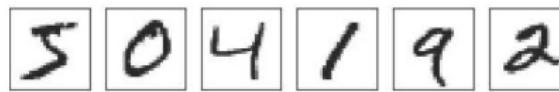
  Applications

  Open Problems

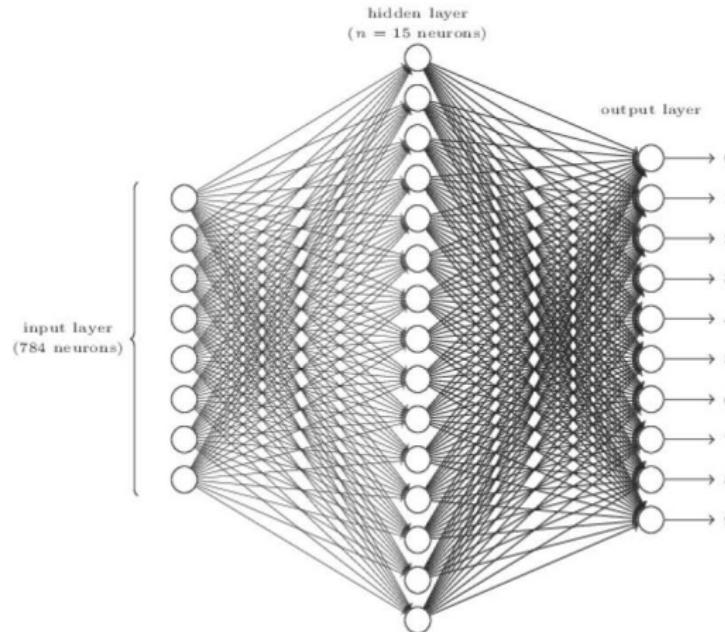
  CNN using Tensorflow

# Classify an Image of a Number

Input:  
(28x28)



Network:



# Outline

Introduction to Deep Learning

  Deep Learning Success

  Why Deep Learning and why now?

TensorFlow

  Introduction to TensorFlow

  TensorFlow Example

  Linear Regression & Gradient Descent

Fundamentals of Deep Learning

  Deep Neural Network

  Artificial Neural Network

Computer Vision & Machine Learning

  Computer Vision in Machine Learning

Artificial Neural Networks (ANN)

  ANN for Computer Vision

  Neural Network for Image Classification

Image Classification using CNN

  Introduction to CNN

  Case Studies of CNN

  Applications

  Open Problems

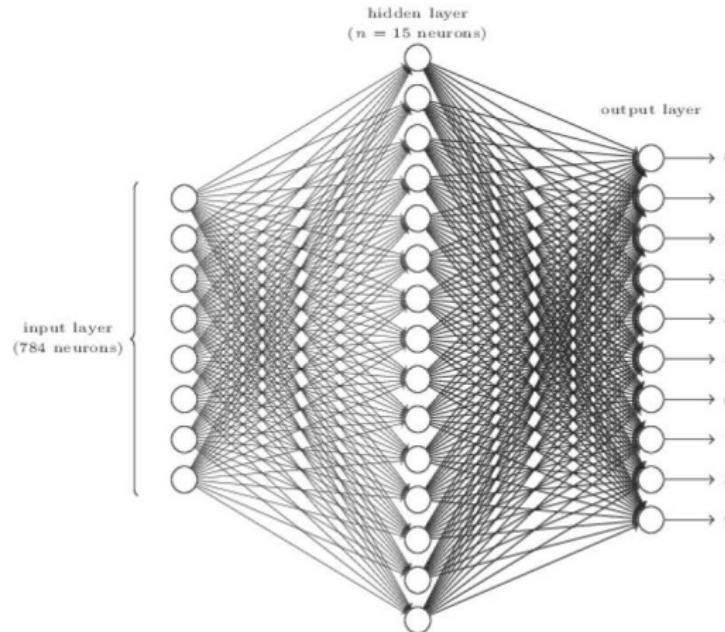
  CNN using Tensorflow

# Classify an Image of a Number

**Input:**  
(28x28)



**Network:**



# Neural Network on MNIST Dataset

- ▶ Creating a simple three layer neural network by using the MNIST dataset that the TensorFlow package provides
- ▶ This MNIST dataset is a set of  $28 \times 28$  pixel grayscale images which represent hand-written digits
- ▶ It has 55,000 training rows, 10,000 testing rows and 5,000 validation rows

## Load Dataset

```
from tensorflow.examples.tutorials.mnist import input_data  
mnist = input_data.read_data_sets("MNIST_data/", one_hot=True)
```



# Neural Network on MNIST Dataset

## Variable Set up

```
# Python optimisation variables
learning_rate = 0.5
epochs = 10
batch_size = 100

# declare the training data placeholders
# input x - for 28 x 28 pixels = 784
x = tf.placeholder(tf.float32, [None, 784])
# now declare the output data placeholder - 10 digits
y = tf.placeholder(tf.float32, [None, 10])
```



# Neural Network on MNIST Dataset

- ▶ There are always L-1 number of weights/bias tensors, where L is the number of layers.

## Weight and Bias Code

```
# now declare the weights connecting the input to the hidden layer  
W1 = tf.Variable(tf.random_normal([784, 300], stddev=0.03), name='W1')  
b1 = tf.Variable(tf.random_normal([300]), name='b1')  
# and the weights connecting the hidden layer to the output layer  
W2 = tf.Variable(tf.random_normal([300, 10], stddev=0.03), name='W2')  
b2 = tf.Variable(tf.random_normal([10]), name='b2')
```

# Neural Network on MNIST Dataset

- ▶ Setup node inputs and activation functions of the hidden layer nodes

## Hidden Layer Node Setup

```
# calculate the output of the hidden layer
hidden_out = tf.add(tf.matmul(x, W1), b1)
hidden_out = tf.nn.relu(hidden_out)
```



# Neural Network on MNIST Dataset

- ▶ Setup the output layer,  $y_-$

## Output Layer Node Setup

```
# now calculate the hidden layer output - in this case, let's use a softmax activated  
# output layer  
y_ = tf.nn.softmax(tf.add(tf.matmul(hidden_out, W2), b2))
```

# Neural Network on MNIST Dataset

- ▶ Optimisation function
  - ▶ Include a cost or loss function for the optimization/backpropagation to work on

## cross entropy cost function

$$J = -\frac{1}{m} \sum_{i=1}^m \sum_{j=1}^n y_j^{(i)} \log(y_{j-}^{(i)}) + (1 - y_j^{(i)}) \log(1 - y_{j-}^{(i)})$$

## Output Layer Node Setup

```
y_clipped = tf.clip_by_value(y_, 1e-10, 0.9999999)
cross_entropy = -tf.reduce_mean(tf.reduce_sum(y * tf.log(y_clipped)
                                             + (1 - y) * tf.log(1 - y_clipped), axis=1))
```

# Neural Network on MNIST Dataset

## Optimizer Algorithm Setup

```
# add an optimiser  
optimiser = tf.train.GradientDescentOptimizer(  
    learning_rate=learning_rate).minimize(cross_entropy)
```

- ▶ This function will then perform the gradient descent and the backpropagation for you.

# Neural Network on MNIST Dataset

- ▶ Setup the variable initialisation operation and an operation to measure the accuracy of our predictions

## Accuracy Code

```
# finally setup the initialisation operator
init_op = tf.global_variables_initializer()

# define an accuracy assessment operation
correct_prediction = tf.equal(tf.argmax(y, 1), tf.argmax(y_, 1))
accuracy = tf.reduce_mean(tf.cast(correct_prediction, tf.float32))
```

# Neural Network on MNIST Dataset

## Setup Code for the Training Process

```
# start the session
with tf.Session() as sess:
    # initialise the variables
    sess.run(init_op)
    total_batch = int(len(mnist.train.labels) / batch_size)
    for epoch in range(epochs):
        avg_cost = 0
        for i in range(total_batch):
            batch_x, batch_y = mnist.train.next_batch(batch_size=batch_size)
            _, c = sess.run([optimiser, cross_entropy],
                           feed_dict={x: batch_x, y: batch_y})
            avg_cost += c / total_batch
        print("Epoch:", (epoch + 1), "cost=", "{:.3f}".format(avg_cost))
    print(sess.run(accuracy, feed_dict={x: mnist.test.images, y: mnist.test.labels}))
```



# Outline

Introduction to Deep Learning

  Deep Learning Success

  Why Deep Learning and why now?

TensorFlow

  Introduction to TensorFlow

  TensorFlow Example

  Linear Regression & Gradient Descent

Fundamentals of Deep Learning

  Deep Neural Network

  Artificial Neural Network

Computer Vision & Machine Learning

  Computer Vision in Machine Learning

Artificial Neural Networks (ANN)

  ANN for Computer Vision

  Neural Network for Image Classification

Image Classification using CNN

**Introduction to CNN**

  Case Studies of CNN

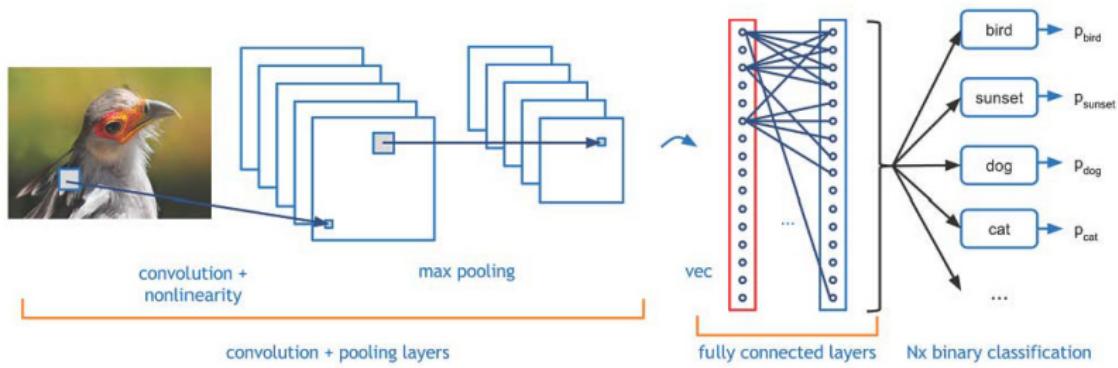
  Applications

  Open Problems

  CNN using Tensorflow

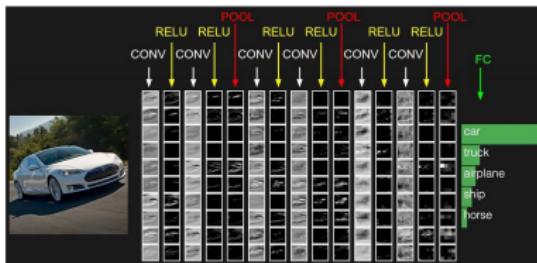
# Convolutional Neural Networks (CNN)

- ▶ Learn a complex representation of visual data using vast amounts of data
- ▶ Inspired by the human visual system
- ▶ Learn multiple layers of transformations, which are applied on top of each other to extract a progressively more sophisticated representation of the input

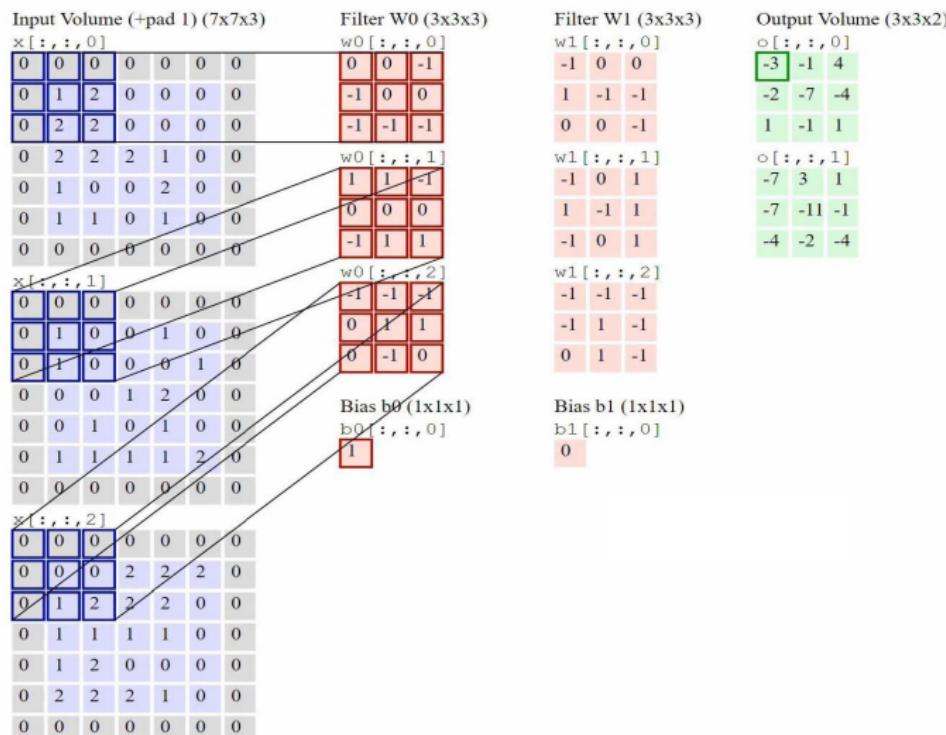


# Convolutional Neural Networks: Layers

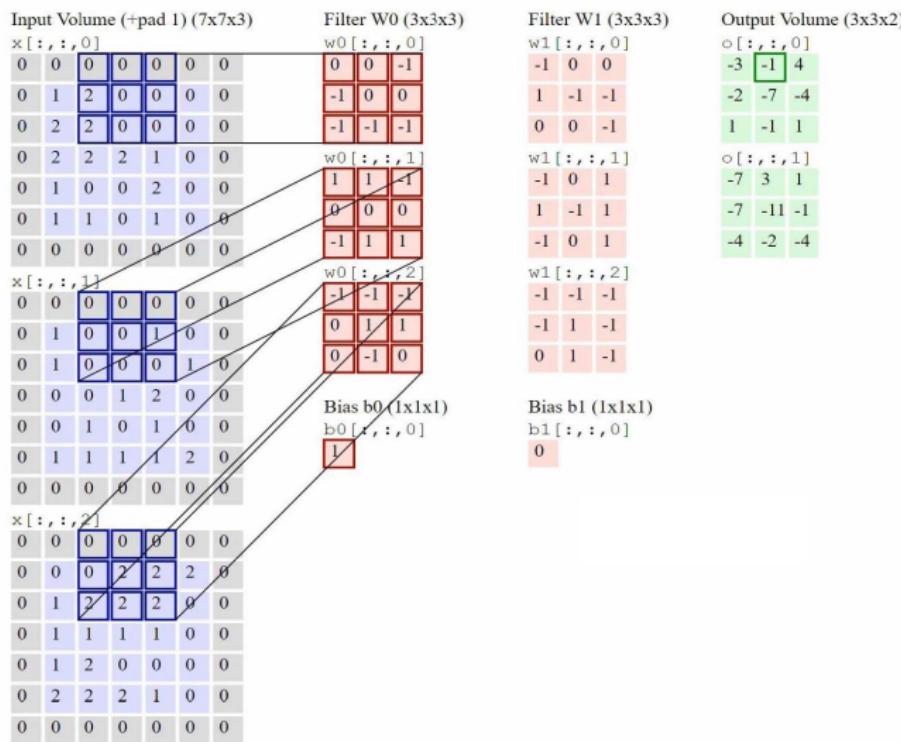
- ▶ **INPUT:**  $[32 \times 32 \times 3]$  will hold the raw pixel values of the image, in this case an image of width 32, height 32, and with three color channels R, G, B
- ▶ **CONV Layer:** will compute the output of neurons that are connected to local regions in the input, each computing a dot product between their weights and a small region they are connected to in the input volume. This may result in volume such as  $[32 \times 32 \times 12]$  if we decided to use 12 filters
- ▶ **RELU Layer:** will apply an elementwise activation function, such as the  $\max(0, x)$  thresholding at zero. This leaves the size of the volume unchanged ( $[32 \times 32 \times 12]$ )
- ▶ **POOL Layer:** will perform a downsampling operation along the spatial dimensions (width, height), resulting in volume such as  $[16 \times 16 \times 12]$
- ▶ **Fully-Connected (FC) Layer:** will compute the class scores, resulting in volume of size  $[1 \times 1 \times 10]$ , where each of the 10 numbers correspond to a class score, such as among the 10 categories of CIFAR-10. As with ordinary Neural Networks and as the name implies, each neuron in this layer will be connected to all the numbers in the previous volume



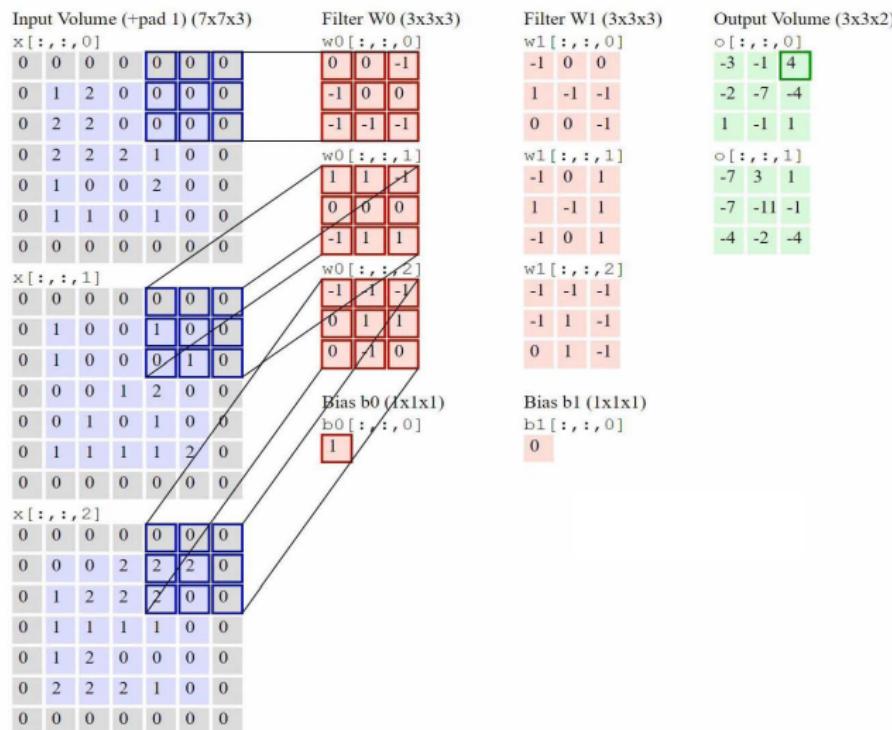
# Convolutional Layer



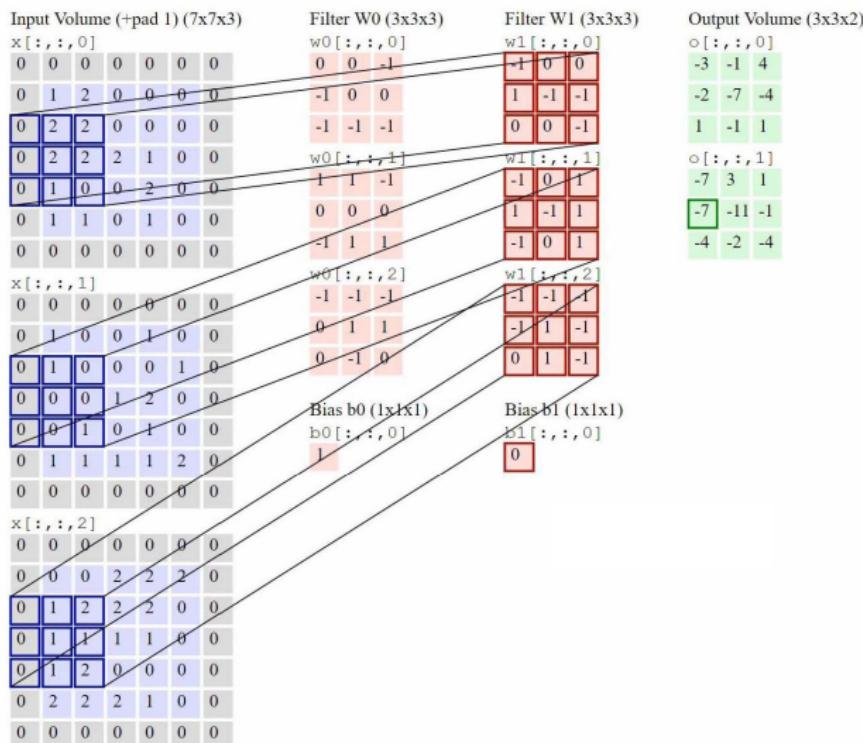
# Convolutional Layer



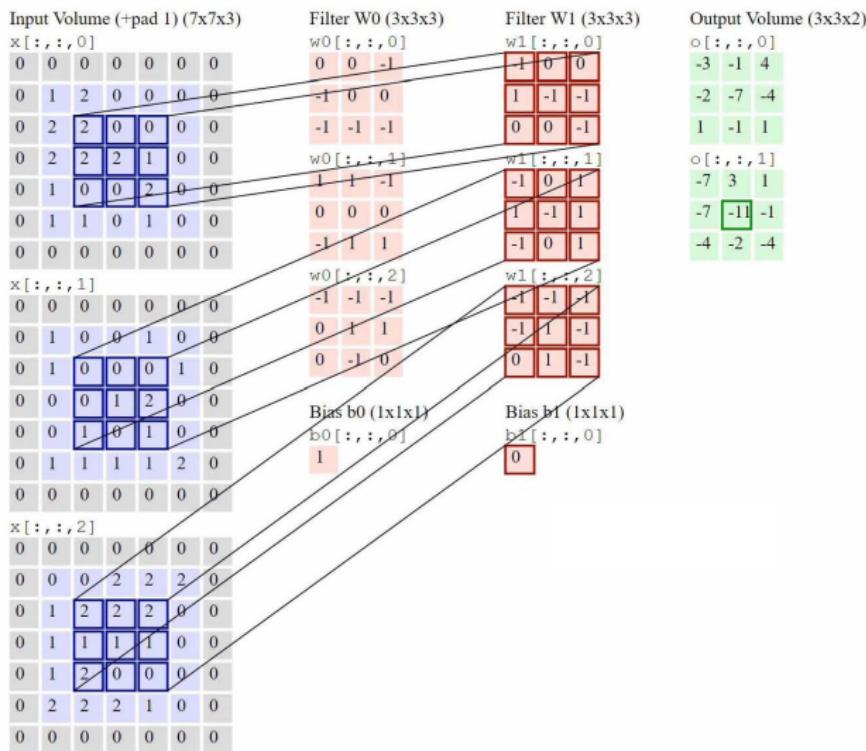
# Convolutional Layer



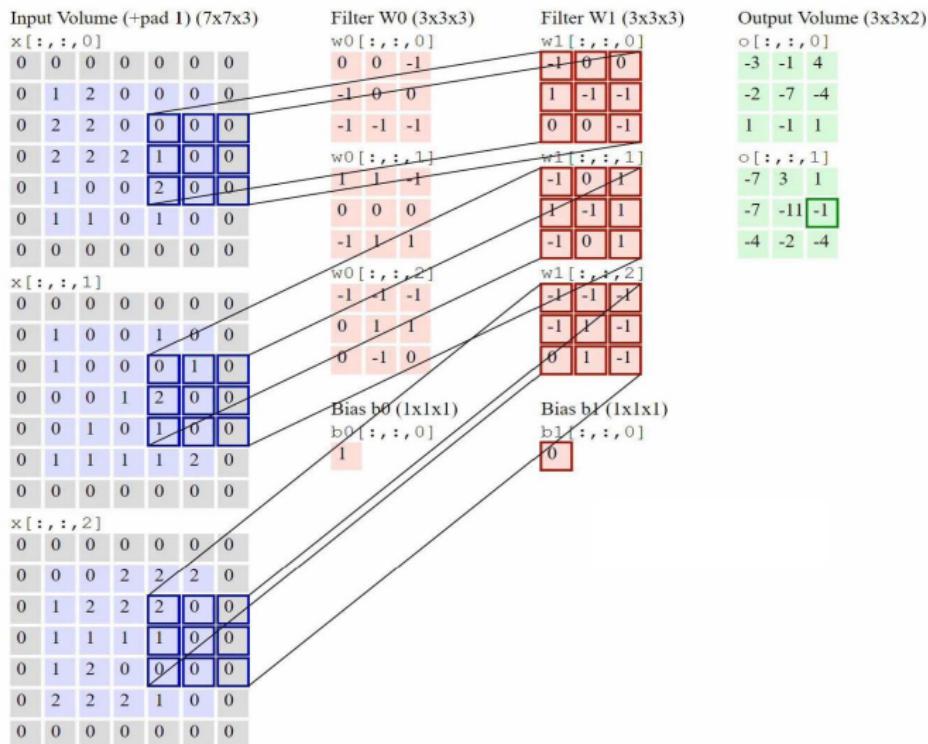
# Convolutional Layer



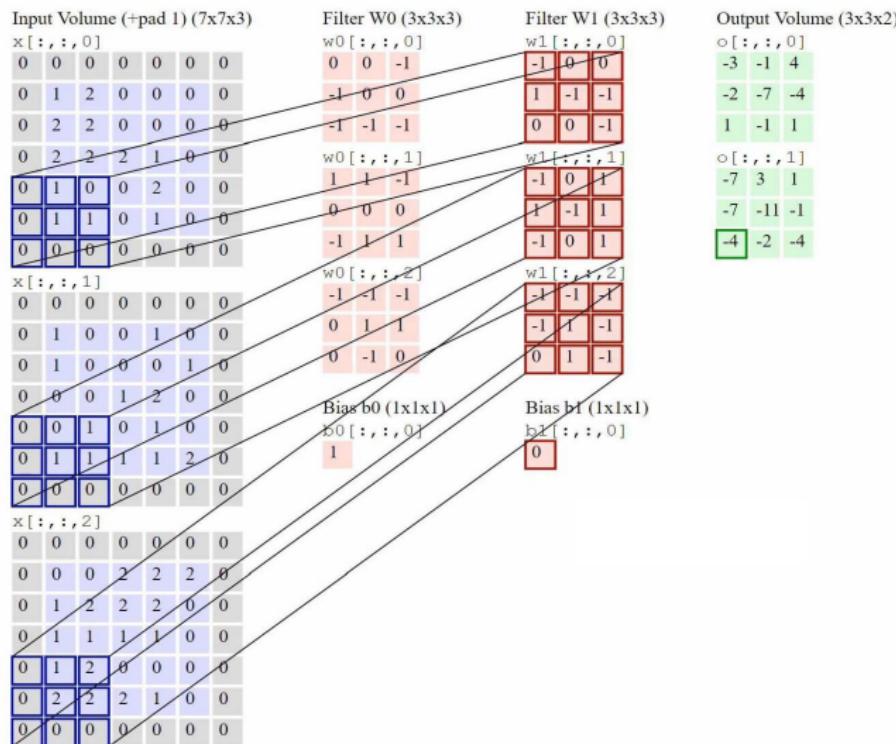
# Convolutional Layer



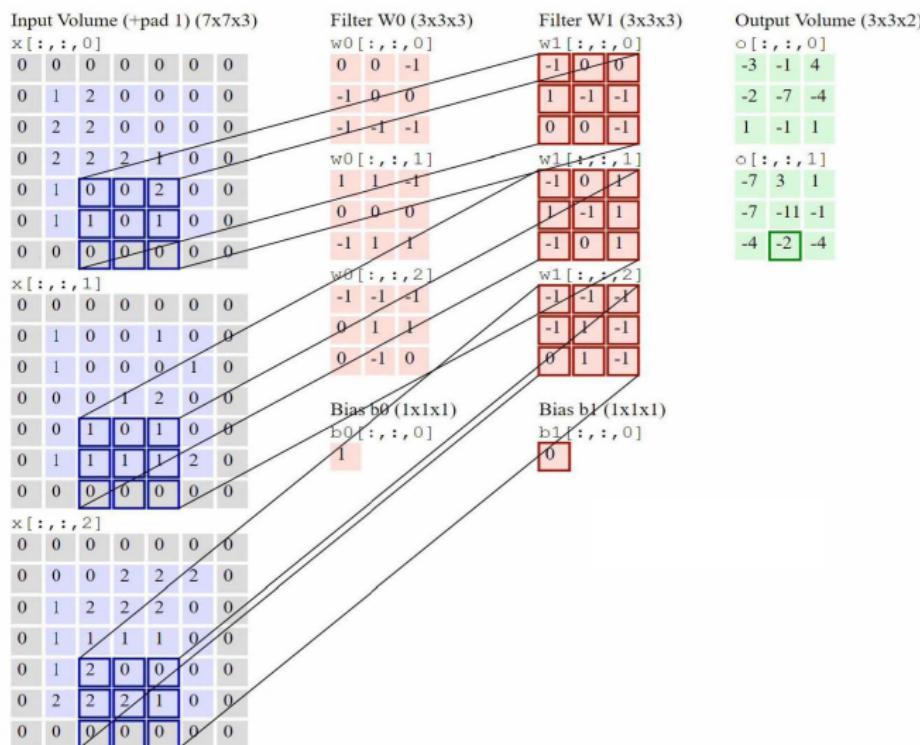
# Convolutional Layer



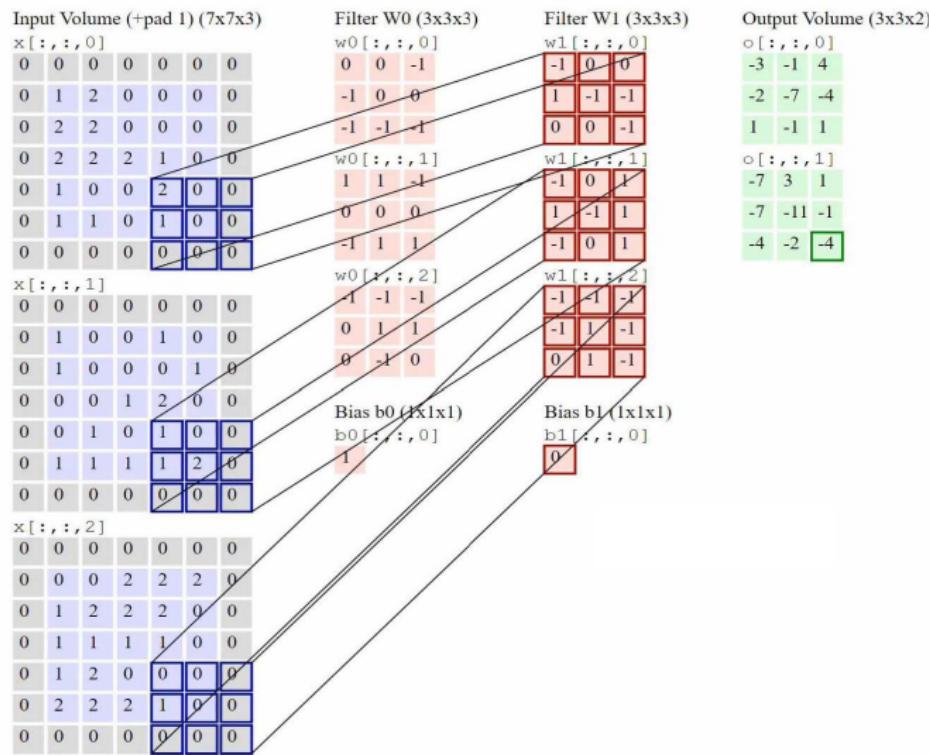
# Convolutional Layer



# Convolutional Layer



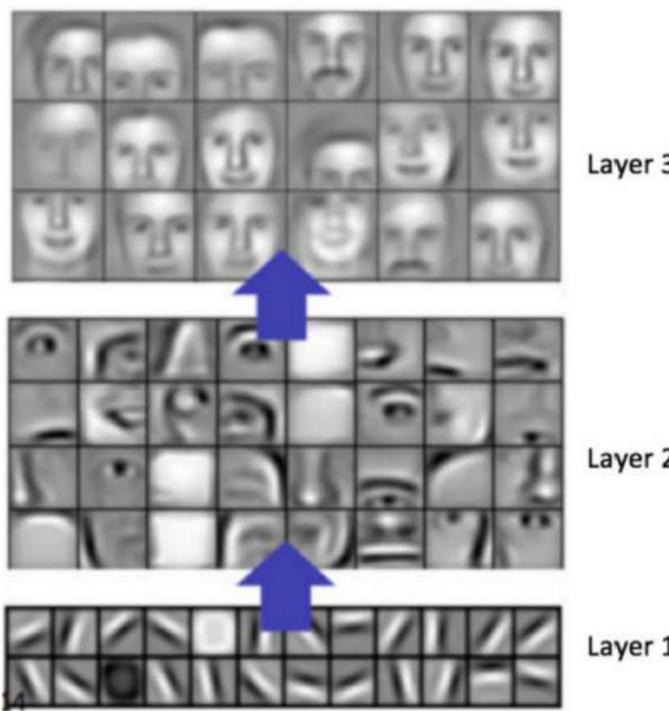
# Convolutional Layer



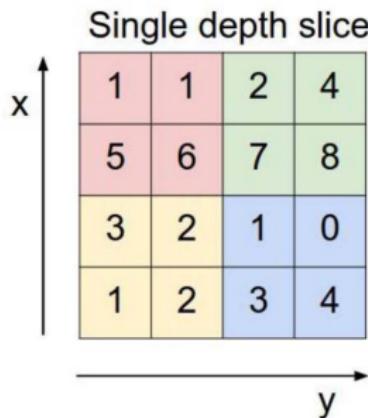
# Convolution

|   | Operation      | Filter  | Convolved Image   |
|---|----------------|---|---|
|  | Identity       | $\begin{bmatrix} 0 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 0 \end{bmatrix}$         |  |
|   |                | $\begin{bmatrix} 1 & 0 & -1 \\ 0 & 0 & 0 \\ -1 & 0 & 1 \end{bmatrix}$       |  |
|   | Edge detection | $\begin{bmatrix} 0 & 1 & 0 \\ 1 & -4 & 1 \\ 0 & 1 & 0 \end{bmatrix}$        |  |
|   |                | $\begin{bmatrix} -1 & -1 & -1 \\ -1 & 8 & -1 \\ -1 & -1 & -1 \end{bmatrix}$ |  |

## Convolution:Representation Learning

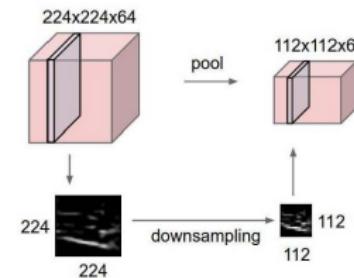


# Pooling Layer



max pool with 2x2 filters  
and stride 2

|   |   |
|---|---|
| 6 | 8 |
| 3 | 4 |



# Outline

Introduction to Deep Learning

Deep Learning Success

Why Deep Learning and why now?

TensorFlow

Introduction to TensorFlow

TensorFlow Example

Linear Regression & Gradient Descent

Fundamentals of Deep Learning

Deep Neural Network

Artificial Neural Network

Computer Vision & Machine Learning

Computer Vision in Machine Learning

Artificial Neural Networks (ANN)

ANN for Computer Vision

Neural Network for Image Classification

Image Classification using CNN

Introduction to CNN

Case Studies of CNN

Applications

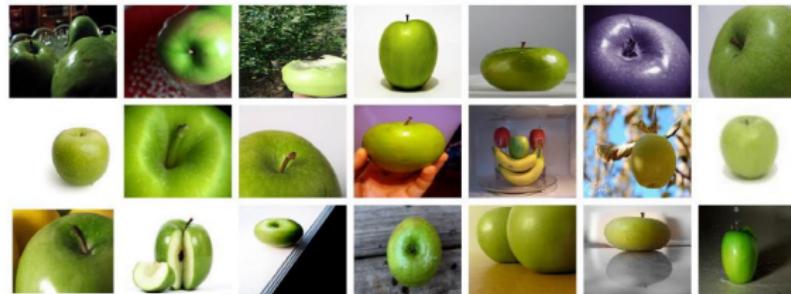
Open Problems

CNN using Tensorflow

# Object Recognition

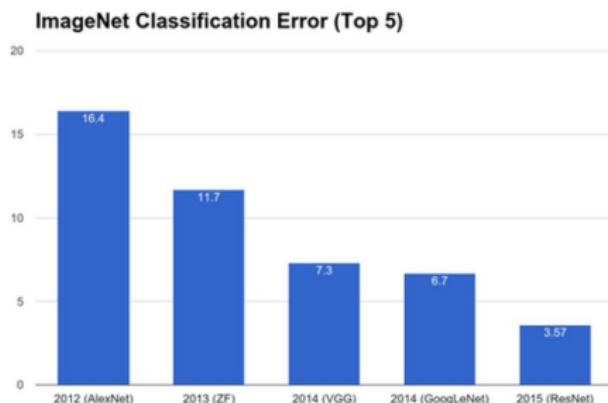
## Case Study: ImageNet

- ▶ **ImageNet:** dataset of 14+ million images (21,841 categories)
  - ▶ Links to images not images
  - ▶ Lets take the high level category of **fruit** as an example:
    - ▶ Total 188,000 images of fruit
    - ▶ There are 1206 Granny Smith apples:



# ImageNet

- ▶ **Competition:** ImageNet Large Scale Visual Recognition Challenge (ILSVRC)
- ▶ Networks
  - ▶ AlexNet (2012)
  - ▶ ZFNet (2013)
  - ▶ VGGNet (2014)
  - ▶ GoogLeNet (2014)
  - ▶ ResNet (2015)
  - ▶ CUIImage (2016)



- **ResNet (2015): 6.7% to 3.57%**
  - More layers = better performance
  - 152 layers
- **CUIImage (2016): 3.57% to 2.99%**
  - Ensemble of 6 models

- **AlexNet (2012): First CNN (15.4%)**

- 8 layers
- 61 million parameters

- **ZFNet (2013): 15.4% to 11.2%**

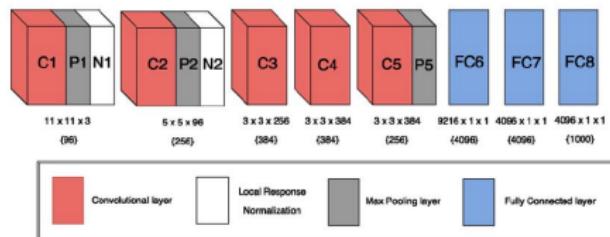
- 8 layers
- More filters. Denser stride.

- **VGGNet (2014): 11.2% to 7.3%**

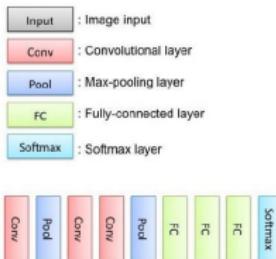
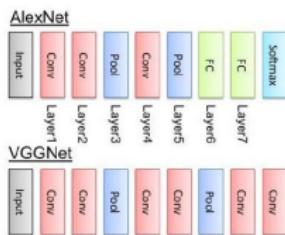
- Beautifully uniform:  
3x3 conv, stride 1, pad 1, 2x2 max pool
- 16 layers
- 138 million parameters

- **GoogLeNet (2014): 11.2% to 6.7%**

- Inception modules
- 22 layers
- 5 million parameters  
(throw away fully connected layers)



- **AlexNet (2012): First CNN (15.4%)**
  - 8 layers
  - 61 million parameters
- **ZFNet (2013): 15.4% to 11.2%**
  - 8 layers
  - More filters. Denser stride.
- **VGGNet (2014): 11.2% to 7.3%**
  - Beautifully uniform:  
3x3 conv, stride 1, pad 1, 2x2 max pool
  - 16 layers
  - 138 million parameters
- **GoogLeNet (2014): 11.2% to 6.7%**
  - Inception modules
  - 22 layers
  - 5 million parameters  
(throw away fully connected layers)

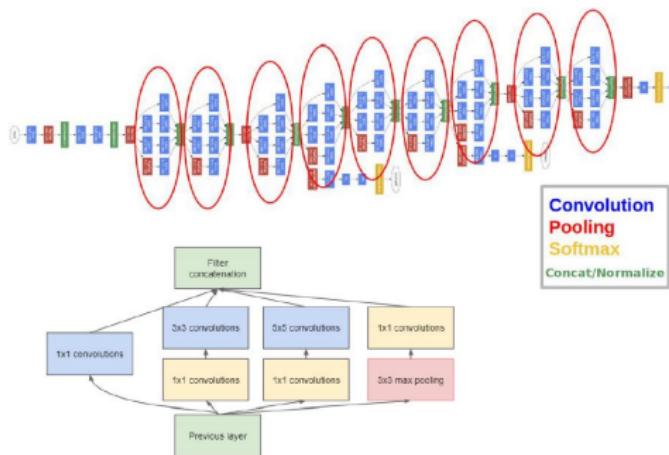


- **ResNet (2015): 6.7% to 3.57%**
  - More layers = better performance
  - 152 layers
- **CUIImage (2016): 3.57% to 2.99%**
  - Ensemble of 6 models

- **AlexNet (2012): First CNN (15.4%)**
  - 8 layers
  - 61 million parameters
- **ZFNet (2013): 15.4% to 11.2%**
  - 8 layers
  - More filters. Denser stride.
- **VGGNet (2014): 11.2% to 7.3%**
  - Beautifully uniform:  
3x3 conv, stride 1, pad 1, 2x2 max pool
  - 16 layers
  - 138 million parameters

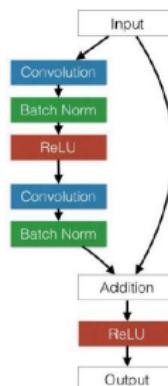
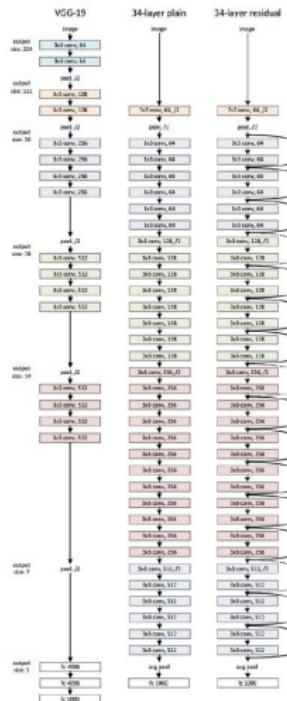
### GoogLeNet (2014): 11.2% to 6.7%

- Inception modules
- 22 layers
- 5 million parameters  
(throw away fully connected layers)



- **ResNet (2015): 6.7% to 3.57%**
  - More layers = better performance
  - 152 layers
- **CUIImage (2016): 3.57% to 2.99%**
  - Ensemble of 6 models

- **AlexNet (2012): First CNN (15.4%)**
  - 8 layers
  - 61 million parameters
- **ZFNet (2013): 15.4% to 11.2%**
  - 8 layers
  - More filters. Denser stride.
- **VGGNet (2014): 11.2% to 7.3%**
  - Beautifully uniform:
    - 3x3 conv, stride 1, pad 1, 2x2 max pool
  - 16 layers
  - 138 million parameters
- **GoogLeNet (2014): 11.2% to 6.7%**
  - Inception modules
  - 22 layers
  - 5 million parameters  
(throw away fully connected layers)



### • VGGNet (2014): 11.2% to 7.3%

- Beautifully uniform:  
3x3 conv, stride 1, pad 1, 2x2 max pool
- 16 layers
- 138 million parameters

### • GoogLeNet (2014): 11.2% to 6.7%

- Inception modules
- 22 layers
- 5 million parameters  
(throw away fully connected layers)

### • ResNet (2015): 6.7% to 3.57%

- More layers = better performance
- 152 layers

### • CUIImage (2016): 3.57% to 2.99%

- Ensemble of 6 models

# Outline

Introduction to Deep Learning

Deep Learning Success

Why Deep Learning and why now?

TensorFlow

Introduction to TensorFlow

TensorFlow Example

Linear Regression & Gradient Descent

Fundamentals of Deep Learning

Deep Neural Network

Artificial Neural Network

Computer Vision & Machine Learning

Computer Vision in Machine Learning

Artificial Neural Networks (ANN)

ANN for Computer Vision

Neural Network for Image Classification

**Image Classification using CNN**

Introduction to CNN

Case Studies of CNN

**Applications**

Open Problems

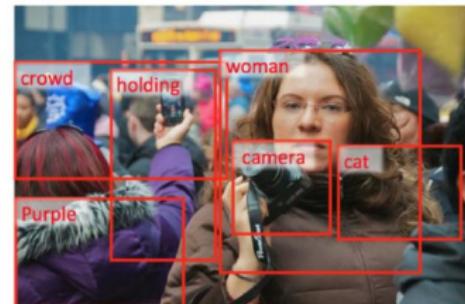
CNN using Tensorflow

# Image Caption Generation



a man sitting on a couch with a dog

a man sitting on a chair with a dog in his lap



# Outline

Introduction to Deep Learning

Deep Learning Success

Why Deep Learning and why now?

TensorFlow

Introduction to TensorFlow

TensorFlow Example

Linear Regression & Gradient Descent

Fundamentals of Deep Learning

Deep Neural Network

Artificial Neural Network

Computer Vision & Machine Learning

Computer Vision in Machine Learning

Artificial Neural Networks (ANN)

ANN for Computer Vision

Neural Network for Image Classification

Image Classification using CNN

Introduction to CNN

Case Studies of CNN

Applications

Open Problems

CNN using Tensorflow

## Fooled by a Little Distortion

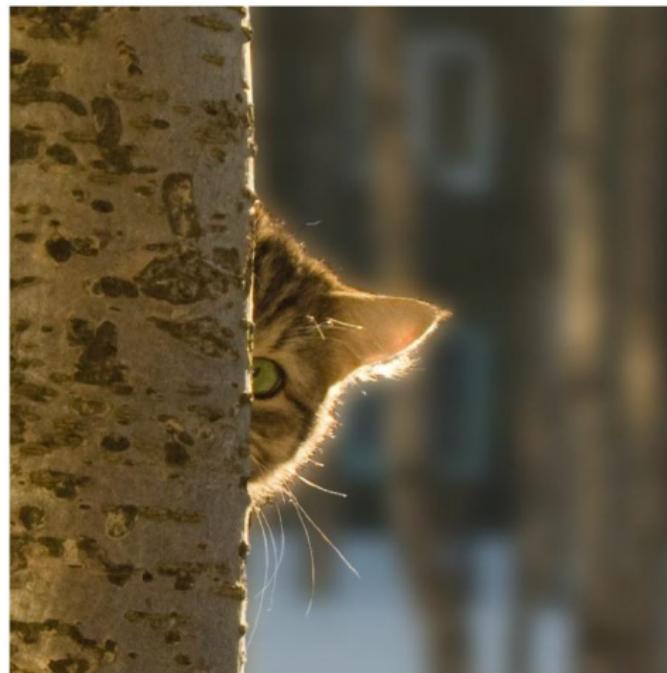


Szegedy et al. "Intriguing properties of neural networks." 2013.

# Object Category Recognition



# Object Category Recognition



## Object Category Recognition



# Outline

Introduction to Deep Learning

  Deep Learning Success

  Why Deep Learning and why now?

TensorFlow

  Introduction to TensorFlow

  TensorFlow Example

  Linear Regression & Gradient Descent

Fundamentals of Deep Learning

  Deep Neural Network

  Artificial Neural Network

Computer Vision & Machine Learning

  Computer Vision in Machine Learning

Artificial Neural Networks (ANN)

  ANN for Computer Vision

  Neural Network for Image Classification

Image Classification using CNN

  Introduction to CNN

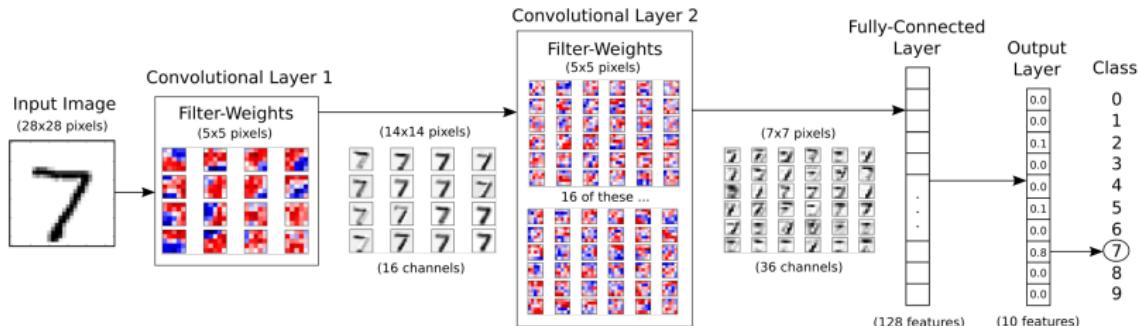
  Case Studies of CNN

  Applications

  Open Problems

CNN using Tensorflow

# CNN Architecture for Image Classification



# CNN on MNIST Dataset

## Load Dataset

```
import tensorflow as tf
import matplotlib.pyplot as plt
#Load dataset
from tensorflow.examples.tutorials.mnist import input_data
mnist = input_data.read_data_sets("/tmp/data/", one_hot = True)
```

## Display of an Image

```
# Get the first images from the test-set.
images = mnist.test.images[0:9]
plt.imshow(images[0].reshape(28,28), cmap='binary')
plt.show()
```



# CNN on MNIST Dataset

## Variable Declaration

```
x = tf.placeholder(tf.float32, shape=[None, 784])
y_ = tf.placeholder(tf.float32, shape=[None, 10])

def weight_variable(shape):
    initial = tf.truncated_normal(shape, stddev=0.1)
    return tf.Variable(initial)

def bias_variable(shape):
    initial = tf.constant(0.1, shape=shape)
    return tf.Variable(initial)
```



# CNN on MNIST Dataset

## TensorFlow Operation for Convolution and Pooling Layer

- ▶ Strides are set to 1 in all dimensions
- ▶ The first and last stride must always be 1, because the first is for the image-number and the last is for the input-channel
- ▶ Strides=[1, 2, 2, 1] would mean that the filter is moved 2 pixels across the x- and y-axis of the image
- ▶ The padding is set to 'SAME' which means the input image is padded with zeroes so the size of the output is the same

```
def conv2d(x, W):
    return tf.nn.conv2d(x, W, strides=[1, 1, 1, 1], padding='SAME')

def max_pool_2x2(x):
    return tf.nn.max_pool(x, ksize=[1, 2, 2, 1],
                          strides=[1, 2, 2, 1], padding='SAME')
```



# CNN on MNIST Dataset

## First Convolution and Pooling Layer

```
#Its weight tensor will have a shape of [5, 5, 1, 32].  
#The first two dimensions are the patch size , the next  
#is the number of input channels ,  
#and the last is the number of output channels.  
W_conv1 = weight_variable([5, 5, 1, 32])  
b_conv1 = bias_variable([32])  
  
#First dim means for all. the second and third dimensions  
#corresponding to image width and height ,  
#and the final dimension corresponding to the number of  
#color channels.  
x_image = tf.reshape(x, [-1, 28, 28, 1])  
  
h_conv1 = tf.nn.relu(conv2d(x_image, W_conv1) + b_conv1)  
h_pool1 = max_pool_2x2(h_conv1)
```

# CNN on MNIST Dataset

## Second Convolution and Pooling Layer

```
W_conv2 = weight_variable([5, 5, 32, 64])  
b_conv2 = bias_variable([64])
```

```
h_conv2 = tf.nn.relu(conv2d(h_pool1, W_conv2) + b_conv2)  
h_pool2 = max_pool_2x2(h_conv2)
```



# CNN on MNIST Dataset

## Fully Connected, Pooling, and Dropout Layer

```
W_fc1 = weight_variable([7 * 7 * 64, 1024])
b_fc1 = bias_variable([1024])

h_pool2_flat = tf.reshape(h_pool2, [-1, 7*7*64])
h_fc1 = tf.nn.relu(tf.matmul(h_pool2_flat, W_fc1) + b_fc1)

keep_prob = tf.placeholder(tf.float32)
h_fc1_drop = tf.nn.dropout(h_fc1, keep_prob)

W_fc2 = weight_variable([1024, 10])
b_fc2 = bias_variable([10])

y_conv = tf.matmul(h_fc1_drop, W_fc2) + b_fc2
```



# CNN on MNIST Dataset

## Cross Entropy

```
cross_entropy = tf.reduce_mean(  
    tf.nn.softmax_cross_entropy_with_logits(labels=y_,  
                                            logits=y_conv))  
train_step = tf.train.AdamOptimizer(1e-4).minimize(cross_entropy)  
correct_prediction = tf.equal(tf.argmax(y_conv, 1),  
                             tf.argmax(y_, 1))  
accuracy = tf.reduce_mean(tf.cast(correct_prediction,  
                                    tf.float32))
```

# CNN on MNIST Dataset

## Session Code

```
with tf.Session() as sess:  
    sess.run(tf.global_variables_initializer())  
    for i in range(2000):  
        batch = mnist.train.next_batch(50)  
        if i % 100 == 0:  
            train_accuracy = accuracy.eval(feed_dict={  
                x: batch[0], y_: batch[1], keep_prob: 1.0})  
            print('step %d, training accuracy %g' % (i,  
                                                train_accuracy))  
        train_step.run(feed_dict={x: batch[0], y_: batch[1],  
                                 keep_prob: 0.5})  
  
    print('test accuracy %g' % accuracy.eval(feed_dict={  
        x: mnist.test.images, y_: mnist.test.labels,  
        keep_prob: 1.0}))
```

## Summary

!For More Details!

<https://www.tensorflow.org>



# Thank You